

Identifiability-Aware Source Apportionment in City-Scale Advection-Diffusion Systems

Ankit Bhardwaj, Lakshmi Subramanian

Department of Computer Science, New York University
bhardwaj.ankit@nyu.edu

Abstract

Source apportionment from sparse urban air-quality sensors is an inverse problem confounded by sensor placement, wind-driven transport, background variation, and noise. Emission inventories restrict the source field to a few candidate groups, but a good predictive fit does not mean those groups can be told apart from the data. Standard receptor models report confident attributions even when they cannot. We cast inventory-based apportionment as a wind-conditioned lagged inverse problem in which each source and temporal-basis coefficient leaves a distinct sensor-time fingerprint. We make identifiability a reported output. After projecting out a low-dimensional background, separability is governed by the projected response matrix $\tilde{H}_\Phi$. Exact recovery needs full column rank, and noise-robust attribution a margin in its singular values, coefficient visibility, background absorption, and pairwise coherence. Our identifiability-aware apportionment (IASA) framework fits nonnegative coefficients, reconstructs activity trajectories, quantifies uncertainty, and, crucially, diagnoses indistinguishable sources and falls back to conservative groupings. Across controlled degeneracies (wind-geometry collapse, source-like backgrounds) IASA flags every non-identifiable case and recovers apportionment to within 0.01 share error. NNLS, CMB, and PMF baselines detect none and misattribute by up to 1.2. On a New Delhi platform (government PM_{2.5} and wind records, regulatory sensor locations, four proxy source groups), IASA reports the finest attribution the declared inventories, transport, background, lag, and noise can support, and no more.

1 Introduction

Urban air pollution is driven by local emissions, meteorological transport, regional background, and sensor noise. For policy the central question is often not only *where* pollution is high but *which* source groups (traffic, industry, brick kilns, population-related activity) produced it, given source inventories, transport models, and sparse measurements.

This is especially hard in sparse sensor networks: a model may fit sensor trajectories accurately yet fail to identify the underlying contributions. The reason is geometric. After wind transport and sparse observation, two distinct source groups can induce nearly identical sensor-time signatures. A strong distant source and a weaker nearby one become observationally equivalent after attenuation, travel time, background correction, and noise. Fine-grained attribution is then not merely uncertain; it is not identifiable from the sensing

system.

We study this in the inventory-based setting. Rather than recovering an arbitrary source field, we assume a finite set of known or proxy source maps that restrict attribution to a low-dimensional, interpretable model. Each source’s time variation is represented as a nonnegative combination of a few temporal basis functions. The unknowns are then source–basis coefficients, from which activity trajectories and contributions are reconstructed. These restrictions alone do not make the problem identifiable: coefficient fingerprints can still be weakly visible, absorbed by background components, or mutually coherent after transport.

The key object is a projected lagged source-response matrix. For a fixed wind realization, transport operator, inventory, sensor layout, and lag window, each source–basis pair induces a finite-horizon sensor-space fingerprint; stacking them gives H_Φ^{lag} . After projecting out a low-dimensional background basis Q , the inverse problem is governed by $\tilde{H}_\Phi = P_Q^\perp H_\Phi^{\text{lag}}$. Its rank, singular values, coefficient visibility, background absorption, pairwise coherence, and ray distance determine which contributions can be separated and which should be reported only after merging or qualification. The goal is therefore not the most detailed contribution vector the inventory permits. It is a conservative resolution, one defensible under the available sensors, wind, background, and noise, and conditional on the declared source model. Estimates are reported alongside identifiability diagnostics and ambiguity certificates.

We make the following contributions:

- We formulate sparse-sensor, inventory-based source apportionment as a lagged wind-conditioned inverse problem over nonnegative source–temporal-basis coefficients, and identify the projected lagged response matrix $\tilde{H}_\Phi = P_Q^\perp H_\Phi^{\text{lag}}$ as the central object governing coefficient separability after transport, sparse observation, lag, and background correction.
- We turn exact and noise-robust identifiability into a concrete reporting protocol built from computable diagnostics of $\tilde{H}_\Phi$ (rank, smallest singular value, condition number, effective rank, coefficient visibility, background absorption, pairwise fingerprint coherence, and ray distance), keeping what is fitted separate from what is diagnosed so that predictive accuracy cannot be mistaken for identifiability.

- We propose an identifiability-aware apportionment (IASA) procedure that fits nonnegative source–basis coefficients, reconstructs source activity trajectories, reports uncertainty, flags weakly visible coefficients, recommends connected report groups when eligible fingerprints are indistinguishable, and adds a per-sensor spatial attribution that back-traces the transport response to upwind regions.
- We instantiate the framework on a New Delhi platform (government PM_{2.5} and wind records, regulatory sensor locations, kernel coordinate-query gridded wind imputation, and four normalized proxy source groups with declared temporal activity bases) and define controlled and observed-data evaluations of conditioning, source ambiguity, background stress, wind diversity, temporal-basis recovery, response error, uncertainty, and residual structure. Across controlled degeneracies, IASA flags every non-identifiable case and recovers apportionment to within 0.01 share error, outperforming the baselines: NNLS, CMB, and PMF.

2 Related Work

Source apportionment has a long history in air-quality science, spanning receptor-oriented methods (chemical mass balance and factor-analytic approaches such as positive matrix factorization (Watson et al. 2002; Paatero and Tapper 1994; Hopke 2016)) and source-oriented models that propagate emissions inventories through transport, with reviews and harmonization efforts documenting their practical sensitivity to source profiles, tracer collinearity, factor interpretation, and inventory assumptions (Viana et al. 2008; Belis et al. 2013, 2019; Mircea et al. 2020; Hopke et al. 2020). This ambiguity is acute in settings such as India, where the same city yields divergent source-contribution estimates and heterogeneously reported categories (Pant and Harrison 2012; Karagulian et al. 2015). Estimating latent quantities from partial observations belongs to inverse problems and data assimilation (Tarantola 2005; Engl, Hanke, and Neubauer 1996; Wang et al. 2023; Carrassi et al. 2018) and, more recently, physics-informed and operator-learning methods (Raissi, Perdikaris, and Karniadakis 2019; Li et al. 2024, 2020; Bhardwaj, Balashankar, and Subramanian 2025). These motivate our transport-structured constrained backend but do not themselves resolve identifiability in a many-to-one sparse-sensing map. Graph, transformer, and sequence models address the same sparse-sensor regime (Li et al. 2017; Yu, Yin, and Zhu 2017; Wu et al. 2019; Chen et al. 2023; Iyer et al. 2022; Bhardwaj et al. 2025), but target forecasting, interpolation, or representation learning rather than attribution to physically meaningful inventories. Finally, our separability question is a form of the observability studied in control theory and system identification (Chen 1984; Ljung 1987; Banks and Kunisch 2012; Colton et al. 1990). Unlike this prior work, we analyze when declared or proxy source groups are *distinguishable* from sparse concentration sensors after wind-conditioned transport, finite lags, background correction, and noise, and report the attribution resolution the sensing system supports. See Appendix K for details.

3 Problem Setup

We formalize source apportionment as a finite-dimensional inverse problem in which known source inventories are transported to a sparse sensor network under a wind-conditioned lagged operator. This defines the observation model, inventory parameterization, response construction, background correction, and the projected response matrix $\tilde{H}_\Phi$ used throughout. Transport, background, and fit details are deferred to Appendix D.

City-scale transport model. Let $u(\mathbf{x}, t)$ be the pollutant concentration field on $\Omega \subset \mathbb{R}^d$, evolving as $\partial u / \partial t = \mathcal{F}_\eta(u, \mathbf{x}, t; w(t)) + s(\mathbf{x}, t) + b(\mathbf{x}, t)$ under meteorological state $w(t)$, local emission s , background b , and transport/dispersion parameters η ; the dynamics may be realized by a plume model, an advection–diffusion solver, or a transport-constrained surrogate.

Sparse observations. Discretize the concentration field on n grid cells, so that $\mathbf{u}_t \in \mathbb{R}^n$. Let $X_{\text{sens}} = \{\mathbf{x}_1, \dots, \mathbf{x}_m\}$ be the fixed sensor locations with $m \ll n$. The observation operator $O : \mathbb{R}^n \rightarrow \mathbb{R}^m$ samples or interpolates the latent field at these locations, $\mathbf{y}_t = O\mathbf{u}_t + \epsilon_t$; stacking over $t = 1, \dots, T$ gives $Y = [\mathbf{y}_1^\top, \dots, \mathbf{y}_T^\top]^\top \in \mathbb{R}^{mT}$. Because m is small relative to the field dimension, many latent fields and source configurations fit the same sensor trajectory.

Inventory-based source parameterization. We assume a set of K known or proxy source maps, $S = [\mathbf{s}_1 \dots \mathbf{s}_K] \in \mathbb{R}^{q \times K}$, where $\mathbf{s}_k \in \mathbb{R}^q$ is the spatial pattern of source group k . Time-varying source activity uses a fixed, predeclared nonnegative temporal dictionary $\Phi = [\phi_1 \dots \phi_B] \in \mathbb{R}_+^{T \times B}$ and coefficients $C = [c_{kb}] \in \mathbb{R}_+^{K \times B}$, so the activity of source k is $\theta_k(t) = \sum_{b=1}^B c_{kb} \phi_b(t) \geq 0$. The basis is a shared dictionary of admissible temporal patterns, not a claim that all sources follow one trajectory: each source has its own nonnegative weights, and source-specific knowledge is imposed by masking inadmissible coefficients. We do not learn Φ from the same sparse observations, since jointly estimating Φ and C would add bilinear non-identifiability beyond the source distinguishability analyzed here. We vectorize C in source-major, basis-minor order as $\mathbf{c} = \text{vec}_{\text{src}}(C) \in \mathbb{R}_+^J$, $J = KB$; the constant-activity model is the special case $B = 1$, $\phi_1(t) = 1$.

Wind-conditioned lagged response. Government wind direction records where wind comes from, whereas transport needs the direction pollution moves; we convert direction WD and speed WS to transport vectors

$$(U_x, V_y) = -WS (\sin \alpha, \cos \alpha), \quad \alpha = WD \pi / 180, \quad (1)$$

and complete them on the response grid with a normalized Gaussian-kernel coordinate-query imputer to obtain a gridded transport field $\hat{\mathbf{w}}_t(\mathbf{x}_g)$; and the same interface supplies synthetic controlled winds and transport ensembles kept separate from inventory scenarios. Emissions released at $t - \ell$ affect sensors at t only after transport by the intervening wind. Let $G_{t, t-\ell}^\partial(w_{t-\ell:t}) \in \mathbb{R}^{n \times q}$ map emissions released on the source grid at $t - \ell$ to the concentration grid at t ; the superscript ∂ denotes an open-boundary operator (mass leaving Ω is removed, not reflected or wrapped). It is realized

by a Gaussian puff approximation evaluated at sensor coordinates that preserves nonnegativity and is mass non-increasing inside the domain, with dispersion, exit bookkeeping, and the pre-fit initial-condition baseline given in Appendix D. For a maximum lag L , $\mathcal{L}_t = \{0, \dots, \min(L, t-1)\}$, the unit-coefficient sensor fingerprint of source k and basis component b at time t is

$$\mathbf{h}_{t,kb}^{\text{lag}} = \sum_{\ell \in \mathcal{L}_t} \phi_b(t-\ell) OG_{t,t-\ell}^{\partial}(w_{t-\ell:t}) \mathbf{s}_k \in \mathbb{R}^m. \quad (2)$$

Stacking over time in source-major, basis-minor order gives $H_{\Phi}^{\text{lag}} \in \mathbb{R}^{mT \times J}$, $J = KB$, each column the finite-horizon sensor-time pattern produced by one unit of a coefficient. The lag window adds no sensor-time rows and is chosen before fitting from physical travel- and residence-time bounds: for adjacent candidates $L, L+\Delta$ with the same rows and columns,

$$\eta_L = \frac{\|H_{\Phi}^{\text{lag}}(L+\Delta) - H_{\Phi}^{\text{lag}}(L)\|_F}{\max\{\|H_{\Phi}^{\text{lag}}(L+\Delta)\|_F, \epsilon\}}, \quad (3)$$

and the primary lag is the smallest candidate with $\eta_L \leq \tau_L$ (default 10^{-3}); the fitted $\mathbf{c}$ never selects L .

Background correction and projection. Observed concentrations include components not explained by the local inventory (regional background, smooth trends, sensor offsets), giving $Y = H_{\Phi}^{\text{lag}} \mathbf{c} + Q\gamma + E$ with a low-dimensional basis $Q \in \mathbb{R}^{mT \times r}$ specified before fitting from source-independent metadata only (timestamps, day labels, sensor identity, and coordinates; effective rank capped at eight, the rank-four for New Delhi observations), so apportionment is not confounded by background correction; columns built from inventories, fingerprints, or fitted signals are excluded and permitted only in labeled stress tests (Appendix D). Real records need not be complete: let $\mathcal{O}$ be the ordered observed sensor-time $\text{PM}_{2.5}$ rows and $M_{\mathcal{O}} \in \{0, 1\}^{N \times mT}$ select them, applied identically to every row-aligned object,

$$Y_{\mathcal{O}} = M_{\mathcal{O}} Y, \quad H_{\Phi, \mathcal{O}}^{\text{lag}} = M_{\mathcal{O}} H_{\Phi}^{\text{lag}}, \quad Q_{\mathcal{O}} = M_{\mathcal{O}} Q, \quad (4)$$

with an alignment mismatch treated as an error; wind may be imputed, $\text{PM}_{2.5}$ is not, and $M_{\mathcal{O}} = I$ for complete controlled data. Dropping the $\mathcal{O}$ subscript, let $P_Q^{\perp} = I - QQ^{\dagger}$ project onto the orthogonal complement of the background space (applied implicitly through a thin SVD of Q); the corrected inverse problem is

$$\tilde{Y} = \tilde{H}_{\Phi} \mathbf{c} + \tilde{E}, \quad \tilde{Y} = P_Q^{\perp} Y, \quad \tilde{H}_{\Phi} = P_Q^{\perp} H_{\Phi}^{\text{lag}}, \quad (5)$$

so each projected fingerprint $\tilde{\mathbf{h}}_{kb} = P_Q^{\perp} \mathbf{h}_{kb}^{\text{lag}}$ encodes the observable effect of one source-basis coefficient after transport, lag, sparse sensing, and background correction. All identifiability diagnostics in this paper are computed from $\tilde{H}_{\Phi}$, not the raw inventory or the unprojected response.

Source activity fit. Scientific knowledge may declare a set $\mathcal{F}_0$ of coefficients identically zero before the response is diagnosed or fitted; with $\mathcal{I} = \{1, \dots, J\} \setminus \mathcal{F}_0$, the fit is computed on $\tilde{H}_{\Phi, \mathcal{I}}$ and zeros are restored afterward (default $\mathcal{F}_0 = \emptyset$). Coefficients are never deleted after fitting merely for being small, which would change the diagnosed problem

after seeing the data. Source-basis coefficients are estimated by nonnegative regularized least squares

$$\hat{\mathbf{c}} = \arg \min_{\mathbf{c} \geq 0} \|\tilde{Y} - \tilde{H}_{\Phi} \mathbf{c}\|_2^2 + \lambda R(\mathbf{c}), \quad (6)$$

with R a ridge or weak inventory-prior penalty, reshaped into activity trajectories $\hat{\theta}_k(t) = \sum_b \hat{c}_{kb} \phi_b(t)$. The regularizer choices, an equivalent joint source/background fit, and an optional constrained end-to-end refinement of wind, dispersion, and coefficients, accepted only if it does not degrade identifiability, are given in Appendix D.

4 Identifiability-Aware Source Apportionment

We now characterize when source activity at a chosen temporal-basis resolution is determined by the projected observations. Throughout, $\tilde{Y} = \tilde{H}_{\Phi} \mathbf{c} + \tilde{E}$ with $\tilde{H}_{\Phi} = P_Q^{\perp} H_{\Phi}^{\text{lag}} \in \mathbb{R}^{N \times J}$, $J = KB$, conditional on the constructed wind field, transport operator, inventories, temporal basis, lag window, and background basis. We use the flattened coefficient index $j = (k, b)$.

Definition 4.1 (Exact identifiability). The source-basis coefficient vector is identifiable from $\tilde{Y}$ if $\tilde{H}_{\Phi} \mathbf{c}_1 = \tilde{H}_{\Phi} \mathbf{c}_2 \implies \mathbf{c}_1 = \mathbf{c}_2$; this is identifiability of the reconstructed source activity trajectories at the selected temporal-basis resolution.

Proposition 4.2 (Identifiability of inventory-based apportionment). *In the noiseless projected model $\tilde{Y} = \tilde{H}_{\Phi} \mathbf{c}$, $\mathbf{c}$ is identifiable uniformly for all $\mathbf{c} \in \mathbb{R}_+^J$ if and only if*

$$\text{rank}(\tilde{H}_{\Phi}) = J. \quad (7)$$

The proof (rank-nullity in both directions, with a nonnegative-feasible counterexample when the rank is deficient) is in Appendix A. We use this full-column-rank condition as our identifiability criterion for the declared coefficient set. Coefficients fixed to zero by scientific knowledge are masked before fitting; the theorem and all diagnostics then apply to the reduced, predeclared column set, and a near-zero fitted estimate never justifies deleting its column after fitting. Rank deficiency may arise from spatial source ambiguity, dependence among temporal modes, or their interaction.

Noise-Robust Identifiability. Let $\sigma_1(\tilde{H}_{\Phi}) \geq \dots \geq \sigma_J(\tilde{H}_{\Phi}) \geq 0$ be the singular values, padded with zeros when $J > N$ or the matrix is rank deficient. Even at full rank, small singular values amplify noise.

Proposition 4.3 (Noise-robust apportionment). *Assume $\tilde{Y} = \tilde{H}_{\Phi} \mathbf{c} + \tilde{E}$ and $\text{rank}(\tilde{H}_{\Phi}) = J$. Let $\hat{\mathbf{c}}_{\text{LS}}$ be the unconstrained least-squares estimate. Then $\|\hat{\mathbf{c}}_{\text{LS}} - \mathbf{c}\|_2 \leq \|\tilde{E}\|_2 / \sigma_J(\tilde{H}_{\Phi})$. If $\hat{\mathbf{c}}$ is the nonnegative least-squares estimate and $\mathbf{c} \geq 0$, then*

$$\|\hat{\mathbf{c}} - \mathbf{c}\|_2 \leq \frac{2\|\tilde{E}\|_2}{\sigma_J(\tilde{H}_{\Phi})}. \quad (8)$$

Both bounds follow from the pseudoinverse identity $\|\tilde{H}_{\Phi}^{\dagger}\|_2 = 1/\sigma_J(\tilde{H}_{\Phi})$ together with feasibility of the true $\mathbf{c}$ in the nonnegative program (Appendix A). Thus $\sigma_J(\tilde{H}_{\Phi})$ measures worst-case robustness: apportionment can be exactly identifiable yet practically unstable when it is small.

Identifiability Diagnostics. Beyond exact rank, IASA reports a fixed panel of diagnostics computed from $\tilde{H}_\Phi$, each with a predeclared, noise- or physics-derived threshold that is never tuned on source recovery (full definitions in Appendix B, threshold values in Table 3, and a summary in Table 1). The *padded minimum singular value* σ_J is the primary score: a rank-deficient response is padded to zero and never reported as stable, and the condition number is $\kappa = \sigma_1/\sigma_J$. The *effective rank* $r_{\text{eff}}(\tau_\sigma)$ counts directions resolvable above the observation-noise level τ_σ ; when it is below J the network cannot support all J coefficients at that noise level. *Coefficient visibility* $v_j = \|\tilde{\mathbf{h}}_j\|_2$ is the post-projection fingerprint magnitude, and the weak set $\mathcal{W} = \{j : v_j \leq \tau_v\}$ flags coefficients below the minimum detectable signal. For eligible (nonweak) pairs, *pairwise coherence* $\rho_{ij} = |\tilde{\mathbf{h}}_i^\top \tilde{\mathbf{h}}_j|/(\|\tilde{\mathbf{h}}_i\|_2 \|\tilde{\mathbf{h}}_j\|_2)$ measures near-proportional signatures, giving the ambiguous-pair set $\mathcal{A} = \{(i, j) : i, j \notin \mathcal{W}, \rho_{ij} > \tau_\rho\}$; an equivalent scale-invariant ray distance $\sqrt{1 - \rho_{ij}^2}$ encodes the same geometry. *Background absorption* $a_j = \|P_Q \mathbf{h}_j^{\text{lag}}\|_2 / \|\mathbf{h}_j^{\text{lag}}\|_2$ is the source-signal fraction lying in the background space, which is why identifiability is assessed after projection. Finally, response error, whether from transport or inventory, acts like additional observation error: it inflates the σ_J bound and is most damaging when $\tilde{H}_\Phi$ is already ill-conditioned, with transport error propagated into uncertainty intervals and inventory error reported only as named robustness scenarios (Appendix E).

Identifiable Source Resolution. When coefficient fingerprints are weak or highly coherent, separate source estimates can mislead. We define the source ambiguity graph $\mathcal{E}_{\text{src}} = \{(k, k') : \exists b, b' ((k, b), (k', b')) \in \mathcal{A}\}$; each edge retains the maximum eligible coefficient coherence, minimum ray distance, and the source–basis pair attaining the trigger. Its deterministically ordered connected components $\mathcal{G} = \{G_1, \dots, G_M\}$ are the recommended report groups. Weak coefficients remain separate flags and create no edges; rank deficiency with no eligible edge produces a global unresolved warning rather than an invented merge. These components are conservative reporting units, not the finest identifiable partition: they can transitively over-merge (edges $A-B$ and $B-C$ group all three even when A and C are distinguishable), so the edge table retains both triggering pairs and their diagnostic values to keep this visible.

The coefficient fit is not silently replaced by a grouped refit. For reporting, group activity and fitted sensor contribution are the member sums $\hat{\theta}_{G_a}(t) = \sum_{k \in G_a} \hat{\theta}_k(t)$ and $\hat{Y}_{G_a} = \sum_{k \in G_a} \sum_b \mathbf{h}_{kb}^{\text{lag}} \hat{c}_{kb}$, while matrix diagnostics remain coefficient level. Table 1 in Appendix C summarizes the scalar diagnostics and the set- and graph-valued reporting outputs.

Per-sensor source footprints. Policy questions are often also spatial. Because the response is linear in $\mathbf{c}$, the fitted solution already localizes each monitor’s signal without a new model: the per-sensor footprint is the puff response of Section 3 read backward from the sensor, and identifiability is inherited rather than created ($\sigma_J(\tilde{H}_\Phi^{(s)}) \leq \sigma_J(\tilde{H}_\Phi)$), so a single sensor is never more identifiable than the pooled

Algorithm 1 Identifiability-Aware Source Apportionment

- 1: **Predeclare** the primary Q , the physical lag-candidate grid, Φ , the inventory version, and the fixed-zero mask $\mathcal{F}_0$, before inspecting any fit.
 - 2: **Wind:** convert WD/WS to transport vectors (Eq. 1), preserving masks/observed values for audit; query the adopted kernel imputer on each response-grid cell for $\hat{\mathbf{w}}_t(\mathbf{x}_g)$ and convert to grid displacement.
 - 3: **Response & lag:** over unmasked source–basis pairs build $\mathbf{h}_{t,kb}^{\text{lag}}$ (Eq. 2); select the smallest L with $\eta_L \leq \tau_L$ from geometry alone (Eq. 3), retaining the full lag sweep.
 - 4: **Align & project:** apply the observed-PM_{2.5} selection identically to Y, H, Q and metadata (Eq. 4) with source-major, basis-minor columns; form $\tilde{Y}, \tilde{H}_\Phi$ (Eq. 5).
 - 5: **Fit:** solve $\hat{\mathbf{c}}$ (Eq. 6) and reconstruct $\hat{\theta}_k(t) = \sum_b \hat{c}_{kb} \phi_b(t)$.
 - 6: **Diagnose:** compute r_{num} , the singular values, $r_{\text{eff}}(\tau_\sigma)$, visibility v_j , absorption a_j , coherence ρ_{ij} , and ray distance d_{ij}^{ray} .
 - 7: **Flag:** weak set $\mathcal{W} = \{j : v_j \leq \tau_v\}$; eligible ambiguous pairs $\mathcal{A} = \{(i, j) : i, j \notin \mathcal{W}, \rho_{ij} > \tau_\rho\}$.
 - 8: **Group:** take deterministic connected components of the source ambiguity graph as conservative report groups (above); weak coefficients are flagged but create no edges, and group activities/contributions are member sums that do not replace the coefficient fit.
 - 9: **Adequacy:** if an externally calibrated noise model exists, run the refitted parametric-bootstrap residual test (below); otherwise label the residual summary uncalibrated.
 - 10: **Report** source estimates, observation/transport uncertainty, inventory robustness scenarios, residual adequacy, diagnostics, and conservative source groups.
-

network). Footprints are therefore reported as spatial overlays of the global fit, with per-sensor source *shares* asserted only at the globally identifiable resolution and aggregated to a report group wherever the pooled diagnostics require one (Appendix F).

IASA Algorithm. We combine the estimator of Section 3 with the diagnostics developed above into an identifiability-aware source apportionment procedure (IASA) that estimates source contributions only at the resolution supported by the projected response matrix. It takes the sparse observations and their row mask, sensors X_{sens} , wind observations and masks, inventories $S \in \mathbb{R}^{q \times K}$, background basis Q , maximum lag L , temporal basis $\Phi \in \mathbb{R}_+^{T \times B}$, and transport response G^∂ , with distinct thresholds $\tau_{\text{num}}, \tau_\sigma, \tau_\rho, \tau_v$ (and the refinement cap τ_ρ^{ref} and wind-drift cap ϵ_w) each predeclared or calibrated from noise or physics and none tuned on source recovery (Appendix G, Table 3); it returns source–basis coefficients and activities at the identifiable grouping, fitted trajectories and residuals, uncertainty intervals, diagnostics, low-visibility flags, and merge recommendations. Algorithm 1 states the procedure.

The step order enforces the method’s central discipline of separating fit from diagnosis: all model choices, thresholds, and the fixed-zero mask are predeclared (line 1), the response is built, projected, and fitted (lines 2–5), and only afterwards are the identifiability diagnostics computed and conservative groups reported (lines 6–8). Because no threshold or column set is chosen after inspecting the fitted coefficients (the lag

itself is selected from response-matrix geometry, not from $\hat{C}$, the reported identifiability cannot be made artificially favorable by post-fit tuning. The objective in Equation 6 is a smooth convex quadratic under a simple nonnegativity constraint, solved with projected FISTA (Beck and Teboulle 2009) (Appendix N).

Model Adequacy, Uncertainty, and Reporting. IASA attaches uncertainty to reported contributions from an active-set/ridge covariance of the fitted coefficients and, when transport ensembles are supplied, from empirical quantiles across ensemble refits. Transport uncertainty and inventory-robustness scenarios are never pooled into one interval. Residual magnitude alone is not calibrated evidence of misspecification. When an externally calibrated observation-noise model is available, IASA computes a background-orthogonal residual statistic whose null distribution (not a plain χ^2 , since $\hat{C}$ is fitted from the same data) is calibrated by a parametric bootstrap (1000 refits, $\alpha = 0.05$). Otherwise it reports raw and projected residual summaries with `calibration_status=uncalibrated` and emits no pass/fail claim. The test is one-sided. A rejection establishes a mismatch with the declared model but not its cause, whereas a non-rejection does not establish inventory completeness, since an omitted source whose signature lies in $\text{span}[H_{\Phi}^{\text{lag}}, Q]$ is absorbed by fitted terms and stays residual-invisible. Prospective network adequacy over historical or simulated wind windows, per-group reliability labeling, the per-sensor contribution decomposition and footprints described above, and acceptance criteria for the optional refinement are detailed in Appendix G.

5 Evaluation

We evaluate on a single New Delhi platform so the controlled and observed studies share one regulatory sensor geometry, one $\text{PM}_{2.5}$ and wind record, and one set of proxy inventories. We ask whether the geometry of $\tilde{H}_{\Phi}$ predicts which coefficient, activity, and grouped-source claims are defensible. A shared base configuration is varied along one controlled axis at a time, avoiding conflation of wind, source geometry, background flexibility, noise, and forward-model error. All runs use the 40×40 resolution, with a 72-hour transport window for the controlled experiments and four consecutive one-week windows for the observed study. Real imputed wind is a controlled transport input when coefficients are synthetically known; the observed- $\text{PM}_{2.5}$ mode is a separate evaluation and is never assigned synthetic source-recovery metrics.

New Delhi Experimental Platform. The platform is built from hourly government $\text{PM}_{2.5}$ and wind records (CPCB 2023, 2025) spanning 1 May 2018–31 October 2020 over a 32-sensor regulatory network. $\text{PM}_{2.5}$ is never imputed, and its valid-value mask defines M_O in Equation 4, with the same ordered sensor-time rows retained in Y , H_{Φ}^{lag} , Q , and metadata. Observed wind is converted (Equation 1) and completed on the response grid by the kernel coordinate-query imputer of Section 3. A learned imputer was trained but not adopted (Section 5.2). Four source groups (brick kilns, industries, population density, traffic) each carry one normalized spatial proxy map and its own admissible temporal-basis components

(Table 4, Appendix H), so fitted coefficients are in normalized proxy units, not physical emission shares. Both forward paths use the open-boundary Gaussian puff operator of Section 3 evaluated at sensors, the matched generator for controlled recovery and the inference operator for observed $\text{PM}_{2.5}$. An auxiliary advection–diffusion simulator supplies a structural forward-model mismatch for the adequacy study (Section 4), and normal backgrounds follow Section 3.

The *controlled* mode preserves this sensor geometry and transport while assigning known nonnegative source-basis coefficients (synthetic or real gridded wind, controlled observation noise, source splits, sensor-layout variants, and optional operator mismatch), so it supports recovery and interval-coverage metrics. The *observed* mode combines real $\text{PM}_{2.5}$, its observation mask, the imputed wind field, and the named normalized inventories; it has no ground-truth source activities and is run over four consecutive one-week windows, each fitted and reported independently. Column-level records, Pusa-monitor averaging and coverage, proxy normalization and crop, temporal-activity assumptions, forward-model baselines, and simulator settings are given in Appendix H.

Metrics and Reporting. Controlled trials, having known coefficients, report coefficient and reconstructed-activity error, grouped activity/contribution error, projected residual, and uncertainty-interval coverage. All trials also report the singular spectrum and its diagnostics (σ_J , numerical and effective rank, condition number, visibility, background absorption, eligible coherence and ray distance), the weak set, triggering pairs, and deterministic report groups (Section 4). They also record lag convergence, the fixed-zero mask, background specification, and ensemble provenance (transport and inventory kept separate). Observed New Delhi runs instead report raw and projected residuals, fitted trajectories, normalized proxy coefficients and contributions, uncertainty intervals, inventory-scenario robustness, weak/ambiguous flags, conservative grouped contributions, and residual calibration status. A reported percentage is a fraction of the fitted inventory-attributed sensor signal, not a physical-emission share.

5.1 Controlled Results

The controlled trials have known source-basis coefficients and therefore recovery and coverage metrics. The controlled geometries are deliberately small (mostly $K = 2$ source groups) so that conditioning, coherence, and background absorption can each be varied in isolation and read directly off $\tilde{H}_{\Phi}$. The observed study then exercises the full four-group, seven-coefficient pipeline on real data. Figure 1 collects the baseline comparison and the seven controlled diagnostics. Full numeric tables are in Appendix I.

Baselines. We compare IASA against three identifiability-blind baselines on the same data: plain NNLS on the projected system (the identifiability layer ablated), chemical mass balance on the unprojected response, and a PMF/NMF receptor factorization. On the wind/geometry collapse (coherence $\rightarrow 1$, $\sigma_J \rightarrow 0$) and the source-like background stress (visibility $\rightarrow 0$, absorption $\rightarrow 1$), IASA raises an identifiability flag in every scenario and seed while none of the baselines do (Figure 1(a); Table 13). Under background stress IASA still recovers ap-




Figure 1: Controlled experiments. **(a)** IASA vs. three identifiability-blind baselines (plain NNLS, chemical mass balance, PMF/NMF) on the wind/geometry collapse and source-like background stress; bars are apportionment share error (log) and the row below marks whether each method raised an identifiability flag (IASA flags both; no baseline does). **(b)** Conditioning, not noise, sets the coefficient-recovery ceiling (close $\sigma_J 9.6/\kappa 2.3$ vs. separated $\sigma_J 3.3/\kappa 8.9$). **(c)** σ_J vs. maximum eligible coherence; colour = wind provider (constant $\rightarrow$ real, dark $\rightarrow$ light), marker = layout (circle regulatory, square random, triangle downwind), size $\propto$ coefficient error; only single/random collapses ($\sigma_J \rightarrow 0$, error 0.57). **(d)** A source-like background drives absorption to 1 and σ_J /visibility to 0. **(e)** Coefficient error grows with the operator-error norm (direction/dispersion axes; speed non-monotonic); structural PDE mismatch gives adequacy rejection 1.0. **(f)** As the lag window grows σ_J rises and κ falls; the rule selects $L=16$ from geometry alone. **(g)** Reconstructed activity error stays below coefficient error at every noise level. **(h)** σ_J tracks the inventory version while recovery stays exact. Values from Appendix I.

portionment to < 0.01 share error whereas the baselines are off by 0.12–1.2. In the collapse the individual split is non-identifiable (coefficient error ≈ 0.57 for the projected fit and unbounded for the raw baselines), yet the baselines report it without warning. A standard method confidently reports a non-identifiable split; IASA reports only the resolution its diagnostics support.

Conditioning and Recovery. We vary source geometry and Gaussian observation noise (0–20% of the maximum clean sensor signal). Both geometries are full rank (numerical rank = effective rank = 2), so the fixed geometry fixes σ_J and κ and coefficient error grows linearly with noise. The better-conditioned *close* geometry recovers proportionally more accurately than the *separated* one (Figure 1(b); Table 5). This is our central claim: response-matrix geometry, not noise alone, sets the recovery ceiling.

Wind Diversity and Sensor Geometry. Matched source, basis, and sensor sets are evaluated across six wind providers (constant, single-direction, diurnal, AR(1), multi-directional, and real gridded New Delhi wind) and three sensor layouts (regulatory, seeded random, downwind-focused), with identical source-basis columns on both sides of every comparison (Figure 1(c); Table 6). Two effects dominate: wind diversity

improves conditioning (multi-directional and real wind give the lowest maximum eligible coherence, while a single steady direction pushes it toward one), and geometry matters as much as wind (under steady wind a seeded random layout collapses σ_J to ≈ 0 and produces the only nonzero coefficient errors in the sweep, whereas the regulatory layout keeps σ_J well away from zero). Separate historical and AR(1) window ensembles both attain full rank with probability 1.0 and no pairwise ambiguity (Table 7).

Background Stress. We compare no background, the normal rank-four basis, a redundant-column basis with the same span, and a labeled source-like stress basis, all declared before fitting. The redundant basis is byte-identical to the primary in every reported quantity, while the source-like *stress* basis drives σ_J and minimum visibility to 0 and absorption to 1 (Figure 1(d); Table 8). A too-flexible background can quietly absorb the source-driven signal and destroy identifiability even when the fitted residual is essentially unchanged. Coefficient error is then misleading: it *falls* in the stress case precisely because the signal has been absorbed.

Transport Error. Parametric perturbations of wind speed, direction, and dispersion within the puff family propagate into attribution. Coefficient error grows with the operator-error

norm along the direction (to 0.86 at 20°) and dispersion (to 0.89 at $2\times$) axes, while the speed axis is non-monotonic, a reminder that a coarse operator-error norm need not order attribution error (Figure 1(e); Table 9). A structural mismatch (open-boundary puff fitted to edge-hold advection–diffusion observations, mismatch norm 0.44) drives mean coefficient error to 0.86 and the refitted-bootstrap adequacy test rejects in 100% of trials, confirming power against forward-model misspecification, not only missing sources.

Lag-Window Selection. Physical travel and residence times define an increasing candidate lag grid. As the window grows the response stabilizes and conditioning improves monotonically (κ falls from 569 at $L=4$ to 5.1 at $L=16$; σ_J rises from 0.08 to 9.07) while numerical rank and report-component count stay fixed and recovery is exact (Figure 1(f); Table 11). The rule (Eq. 3) picks the smallest lag whose relative Frobenius change falls to $\tau_L = 10^{-3}$ (here $L=16$) from geometry alone, never from fitted coefficients.

Temporal-Basis Recovery. Reconstructing known diurnal, intermittent, and day/night source–basis coefficients at increasing noise separates coefficient error from reconstructed-activity error. The activity-trajectory error stays consistently below the coefficient error (0.58 vs. 0.90 at 20% noise), at a near-constant ratio ≈ 0.57 (Figure 1(g); Table 12), because projecting noisy coefficients through the temporal bases damps isotropic estimation noise more than the coherent signal.

Inventory Robustness. Perturbing inventory locations, spatial scales, category assignments, and map versions (each a separate scenario, never pooled with transport error) recovers the known coefficients exactly, but the identifiability geometry moves with the inventory version. Rescaling or substituting the map nearly doubles σ_J (to 47.0 and 46.0 vs. 21.6 baseline) while a symmetric category swap leaves it unchanged (Figure 1(h); Table 10). The attribution *target* thus changes with the assumed inventory even when recovery is exact. Three further controlled experiments not shown here (coherence and grouped reporting, missing-source adequacy, and per-sensor footprints) are reported in Appendix J.

5.2 Observed New Delhi

New Delhi Wind-Imputation Validation. The learned wind model did not beat a non-spatial city-mean baseline on held-out station–time vectors, so the spatially resolved kernel imputer is used throughout (the city mean lacks the spatial resolution the response operator requires). Dense wind truth is unavailable, so gridded-field accuracy is assessed only through the controlled synthetic-wind experiments, where the true wind is known.

New Delhi Identifiability and Report Groups. The observed study runs as four consecutive one-week windows (2018-05-01 through 2018-05-28), each a self-contained fit; results are not aggregated across windows. This window is chosen for its near-complete early-record coverage. It is a demonstration window, the full 2018–2020 record was computationally infeasible. Every week is full numerical rank (7 free coefficients after the fixed-zero mask), finite-conditioned, with low maximum eligible coherence (0.32–0.50), an empty weak set, no cross-source ambiguous pairs, and four singleton




Figure 2: Observed New Delhi. (a) Weeks 1–4: every week is full rank; bars σ_1/σ_J (left, log) and max eligible coherence (red, right axis 0–1) (Table 14). (b) proxy apportionment as a fraction of fitted sensor signal (not physical emissions), swinging population- to brick-kiln-dominated (Table 15).

report components (Figure 2(a); Table 14). Each source group is therefore individually identifiable in every week and no conservative merge is required.

New Delhi Proxy Apportionment and Uncertainty. Figure 2(b) reports each week’s proxy apportionment as a fraction of the fitted inventory-attributed sensor signal (the sum over groups of the L_1 fitted per-group sensor-signal magnitude). These are explicitly *not* physical-emission shares. The mixture swings substantially, population-dominated in weeks 1–3 and brick-kiln-dominated in week 4, which is why the windows are reported separately rather than as one aggregate estimate. The fit uses only the 79–85% of rows with a valid $\text{PM}_{2.5}$ measurement (no imputation), subtracting each week’s zero-source transported kriged initial-condition baseline (Pusa monitors averaged) under the gridded kernel wind field. The week-4 brick-kiln dominance (Figure 2(b)) fits the end-of-season firing push before the monsoon.

6 Discussion and Conclusion

We framed sparse-sensor apportionment as an identifiable-resolution problem: inventories make attribution meaningful but need not make source groups separable. Separability is governed by the projected response matrix $\tilde{H}_\Phi$. Exact identifiability requires $\text{rank}(\tilde{H}_\Phi) = J$, and robust attribution its singular values, visibility, absorption, and coherence. Prediction accuracy alone is therefore not evidence of identifiability. IASA reports contributions only where this geometry supports them and qualifies them otherwise, and judges sensor, wind, and background design by fingerprint conditioning and coherence, not coverage. All claims are conditional on the declared wind, transport, basis, inventory, and background. The basis bounds recoverable activity, aligned omitted sources stay residual-invisible, and transport uncertainty (intervals) stays separate from inventory scenarios (unpooled). Richer meteorology, spatial diagnostics, and joint basis, transport, and inventory estimation are natural next steps.

Acknowledgements

Ankit Bhardwaj and Lakshminarayanan Subramanian were supported by the NSF Grant (Award Number 2335773) titled

“EAGER: Scalable Climate Modeling using Message-Passing Recurrent Neural Networks”. Lakshminarayanan Subramanian was also funded in part by the NSF Grant (award number OAC-2004572) titled “A Data-informed Framework for the Representation of Sub-grid Scale Gravity Waves to Improve Climate Prediction”.

References

- Banks, H. T.; and Kunisch, K. 2012. *Estimation techniques for distributed parameter systems*. Springer Science & Business Media.
- Beck, A.; and Teboulle, M. 2009. A fast iterative shrinkage-thresholding algorithm for linear inverse problems. *SIAM Journal on Imaging Sciences*, 2(1): 183–202.
- Belis, C.; Favez, O.; Mircea, M.; Diapouli, E.; Manousakas, M.; Vratolis, S.; Gilardoni, S.; Paglione, M.; Decesari, S.; Mocnik, G.; et al. 2019. European guide on air pollution source apportionment with receptor models—Revised version 2019. *Publications Office, LU*.
- Belis, C.; Karagulian, F.; Larsen, B. R.; and Hopke, P. 2013. Critical review and meta-analysis of ambient particulate matter source apportionment using receptor models in Europe. *Atmospheric Environment*, 69: 94–108.
- Bhardwaj, A.; Balashankar, A.; Iyer, S.; Soans, N.; Sudarshan, A.; Pande, R.; and Subramanian, L. 2025. Comprehensive monitoring of air pollution hotspots using sparse sensor networks. *ACM Journal on Computing and Sustainable Societies*, 3(4): 1–36.
- Bhardwaj, A.; Balashankar, A.; and Subramanian, L. 2025. FieldFormer: Physics-Informed Transformers for Spatio-Temporal Field Reconstruction from Sparse Sensors. *arXiv preprint arXiv:2510.03589*.
- Carrassi, A.; Bocquet, M.; Bertino, L.; and Evensen, G. 2018. Data assimilation in the geosciences: An overview of methods, issues, and perspectives. *Wiley Interdisciplinary Reviews: Climate Change*, 9(5): e535.
- Center for International Earth Science Information Network (CIESIN), Columbia University. 2018. Gridded Population of the World, Version 4 (GPWv4): Population Density, Revision 11.
- Chen, C.-T. 1984. *Linear system theory and design*, volume 301. Holt, Rinehart and Winston New York.
- Chen, L.; Xu, J.; Wu, B.; and Huang, J. 2023. Group-aware graph neural network for nationwide city air quality forecasting. *ACM Transactions on Knowledge Discovery from Data*, 18(3): 1–20.
- Colton, D. L.; Ewing, R. E.; Rundell, W.; et al. 1990. *Inverse problems in partial differential equations*, volume 42. Siam.
- CPCB. 2023. Air Quality Data. <https://cpcb.nic.in/>.
- CPCB. 2025. CPCB Data Portal. <https://app.cpcbcr.com/ccr/#/caaqm-dashboard-all/caaqm-landing/caaqm-comparison-data>.
- Engl, H. W.; Hanke, M.; and Neubauer, A. 1996. *Regularization of inverse problems*, volume 375. Springer Science & Business Media.
- Google. 2024. Google Maps. <https://maps.google.com>. Accessed: 2026-03-30.
- Guttikunda, S. K.; and Calori, G. 2013. A GIS based emissions inventory at 1 km × 1 km spatial resolution for air pollution analysis in Delhi, India. *Atmospheric Environment*, 67: 101–111.
- Hopke, P. K. 2016. Review of receptor modeling methods for source apportionment. *Journal of the Air & Waste Management Association*, 66(3): 237–259.
- Hopke, P. K.; Dai, Q.; Li, L.; and Feng, Y. 2020. Global review of recent source apportionments for airborne particulate matter. *Science of The Total Environment*, 740: 140091.
- Iyer, S. R.; Balashankar, A.; Aeberhard, W. H.; Bhattacharyya, S.; Rusconi, G.; Jose, L.; Soans, N.; Sudarshan, A.; Pande, R.; and Subramanian, L. 2022. Modeling fine-grained spatio-temporal pollution maps with low-cost sensors. *npj Climate and Atmospheric Science*, 5(1): 76.
- Karagulian, F.; Belis, C. A.; Dora, C. F. C.; Prüss-Ustün, A. M.; Bonjour, S.; Adair-Rohani, H.; and Amann, M. 2015. Contributions to cities’ ambient particulate matter (PM): A systematic review of local source contributions at global level. *Atmospheric environment*, 120: 475–483.
- Li, Y.; Yu, R.; Shahabi, C.; and Liu, Y. 2017. Diffusion convolutional recurrent neural network: Data-driven traffic forecasting. *arXiv preprint arXiv:1707.01926*.
- Li, Z.; Kovachki, N.; Azizzadenesheli, K.; Liu, B.; Bhattacharya, K.; Stuart, A.; and Anandkumar, A. 2020. Fourier neural operator for parametric partial differential equations. *arXiv preprint arXiv:2010.08895*.
- Li, Z.; Zheng, H.; Kovachki, N.; Jin, D.; Chen, H.; Liu, B.; Azizzadenesheli, K.; and Anandkumar, A. 2024. Physics-informed neural operator for learning partial differential equations. *ACM/IMS Journal of Data Science*, 1(3): 1–27.
- Ljung, L. 1987. *System Identification: Theory for the User*. Prentice Hall.
- Mircea, M.; Calori, G.; Pirovano, G.; Belis, C.; et al. 2020. European guide on air pollution source apportionment for particulate matter with source oriented models and their combined use with receptor models. *Publications Office of the European Union, LU*.
- Paatero, P.; and Tapper, U. 1994. Positive matrix factorization: A non-negative factor model with optimal utilization of error estimates of data values. *Environmetrics*, 5(2): 111–126.
- Pant, P.; and Harrison, R. M. 2012. Critical review of receptor modelling for particulate matter: a case study of India. *Atmospheric Environment*, 49: 1–12.
- Raissi, M.; Perdikaris, P.; and Karniadakis, G. E. 2019. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational physics*, 378: 686–707.
- Seber, G. A. F.; and Lee, A. J. 2003. *Linear Regression Analysis*. Hoboken, NJ: Wiley, 2nd edition.
- Tarantola, A. 2005. *Inverse problem theory and methods for model parameter estimation*. SIAM.

Viana, M.; Kuhlbusch, T. A.; Querol, X.; Alastuey, A.; Harrison, R. M.; Hopke, P. K.; Winiwarter, W.; Vallius, M.; Szidat, S.; Prévôt, A. S.; et al. 2008. Source apportionment of particulate matter in Europe: a review of methods and results. *Journal of aerosol science*, 39(10): 827–849.

Wang, B.; Sun, Z.; Jiang, X.; Zeng, J.; and Liu, R. 2023. Kalman filter and its application in data assimilation. *Atmosphere*, 14(8): 1319.

Watson, J. G.; Zhu, T.; Chow, J. C.; Engelbrecht, J.; Fujita, E. M.; and Wilson, W. E. 2002. Receptor modeling application framework for particle source apportionment. *Chemosphere*, 49(9): 1093–1136.

Wu, Z.; Pan, S.; Long, G.; Jiang, J.; and Zhang, C. 2019. Graph wavenet for deep spatial-temporal graph modeling. *arXiv preprint arXiv:1906.00121*.

Yu, B.; Yin, H.; and Zhu, Z. 2017. Spatio-temporal graph convolutional networks: A deep learning framework for traffic forecasting. *arXiv preprint arXiv:1709.04875*.

A Detailed Identifiability Proofs

Proof of Proposition 4.2. Let $A = \tilde{H}_\Phi \in \mathbb{R}^{N \times J}$. If $\text{rank}(A) = J$, then the rank-nullity theorem gives $\dim \ker(A) = J - \text{rank}(A) = 0$. Hence $\ker(A) = \{\mathbf{0}\}$. If two feasible coefficient vectors $\mathbf{c}_1, \mathbf{c}_2 \in \mathbb{R}_+^J$ produce the same projected observation,

$$A\mathbf{c}_1 = A\mathbf{c}_2,$$

then $A(\mathbf{c}_1 - \mathbf{c}_2) = \mathbf{0}$. Since the only vector in the kernel is $\mathbf{0}$, we have $\mathbf{c}_1 - \mathbf{c}_2 = \mathbf{0}$, and therefore $\mathbf{c}_1 = \mathbf{c}_2$. Thus the coefficient vector is identifiable over the nonnegative feasible set.

Conversely, suppose $\text{rank}(A) < J$. By rank-nullity, $\dim \ker(A) > 0$, so there exists a nonzero vector $\Delta\mathbf{c} \in \ker(A)$. Choose any $\beta > \|\Delta\mathbf{c}\|_\infty$ and set $\mathbf{c} = \beta\mathbf{1}$. Then $\mathbf{c} \in \mathbb{R}_+^J$, and $\mathbf{c} + \Delta\mathbf{c} \in \mathbb{R}_+^J$ because every component satisfies

$$\beta + \Delta c_j \geq \beta - \|\Delta\mathbf{c}\|_\infty > 0.$$

The two vectors are distinct because $\Delta\mathbf{c} \neq \mathbf{0}$, but they produce the same projected observation:

$$A(\mathbf{c} + \Delta\mathbf{c}) = A\mathbf{c} + A\Delta\mathbf{c} = A\mathbf{c}.$$

Therefore uniform identifiability over $\mathbb{R}_+^J$ fails whenever $\text{rank}(A) < J$. This proves the equivalence.

Proof of Proposition 4.3. Again let $A = \tilde{H}_\Phi$, and assume A has full column rank J . The projected observation model is

$$\tilde{Y} = A\mathbf{c} + \tilde{E}.$$

For unconstrained least squares,

$$\hat{\mathbf{c}}_{\text{LS}} = \arg \min_{\mathbf{z}} \|\tilde{Y} - A\mathbf{z}\|_2^2 = A^\dagger \tilde{Y},$$

where $A^\dagger$ is the Moore–Penrose pseudoinverse. Substituting the data model gives

$$\hat{\mathbf{c}}_{\text{LS}} = A^\dagger (A\mathbf{c} + \tilde{E}) = \mathbf{c} + A^\dagger \tilde{E},$$

because $A^\dagger A = I_J$ for a full-column-rank matrix. Hence

$$\|\hat{\mathbf{c}}_{\text{LS}} - \mathbf{c}\|_2 = \|A^\dagger \tilde{E}\|_2 \leq \|A^\dagger\|_2 \|\tilde{E}\|_2.$$

The singular values of $A^\dagger$ are the reciprocals of the nonzero singular values of A , so $\|A^\dagger\|_2 = 1/\sigma_J(A)$. Therefore

$$\|\hat{\mathbf{c}}_{\text{LS}} - \mathbf{c}\|_2 \leq \frac{\|\tilde{E}\|_2}{\sigma_J(A)}.$$

For nonnegative least squares, let $\hat{\mathbf{c}}$ minimize $\|\tilde{Y} - A\mathbf{z}\|_2^2$ over $\mathbf{z} \geq \mathbf{0}$. Since the true $\mathbf{c} \geq \mathbf{0}$, it is feasible. Optimality of $\hat{\mathbf{c}}$ therefore implies

$$\|\tilde{Y} - A\hat{\mathbf{c}}\|_2 \leq \|\tilde{Y} - A\mathbf{c}\|_2 = \|\tilde{E}\|_2.$$

Let $\Delta = \hat{\mathbf{c}} - \mathbf{c}$. Using $\tilde{Y} = A\mathbf{c} + \tilde{E}$,

$$\tilde{Y} - A\hat{\mathbf{c}} = \tilde{E} - A\Delta.$$

Thus

$$\begin{aligned} \|A\Delta\|_2 &= \|\tilde{E} - (\tilde{E} - A\Delta)\|_2 \\ &\leq \|\tilde{E}\|_2 + \|\tilde{E} - A\Delta\|_2 \leq 2\|\tilde{E}\|_2. \end{aligned}$$

Full column rank gives the lower singular-value inequality

$$\|A\Delta\|_2 \geq \sigma_J(A) \|\Delta\|_2.$$

Combining the two inequalities yields

$$\|\hat{\mathbf{c}} - \mathbf{c}\|_2 = \|\Delta\|_2 \leq \frac{2\|\tilde{E}\|_2}{\sigma_J(A)}.$$

This proves the stated NNLS robustness bound.

Ray distance for nonweak fingerprints (Equation 9). Fix nonweak projected fingerprints $\tilde{\mathbf{h}}_i, \tilde{\mathbf{h}}_j$, so $\tilde{\mathbf{h}}_j \neq \mathbf{0}$. The ray distance minimizes

$$g(\alpha) = \frac{\|\tilde{\mathbf{h}}_i - \alpha \tilde{\mathbf{h}}_j\|_2^2}{\|\tilde{\mathbf{h}}_j\|_2^2}$$

over $\alpha \in \mathbb{R}$. This is strictly convex in α , and $g'(\alpha) = 0$ gives $-2\tilde{\mathbf{h}}_j^\top (\tilde{\mathbf{h}}_i - \alpha \tilde{\mathbf{h}}_j) = 0$, i.e.

$$\alpha^* = \frac{\tilde{\mathbf{h}}_i^\top \tilde{\mathbf{h}}_j}{\|\tilde{\mathbf{h}}_j\|_2^2}.$$

Substituting,

$$\|\tilde{\mathbf{h}}_i - \alpha^* \tilde{\mathbf{h}}_j\|_2^2 = \|\tilde{\mathbf{h}}_i\|_2^2 - \frac{(\tilde{\mathbf{h}}_i^\top \tilde{\mathbf{h}}_j)^2}{\|\tilde{\mathbf{h}}_j\|_2^2},$$

and dividing by $\|\tilde{\mathbf{h}}_i\|_2^2$ yields $(d_{ij}^{\text{ray}})^2 = 1 - \rho_{ij}^2$, so $d_{ij}^{\text{ray}} = \sqrt{1 - \rho_{ij}^2}$. Minimizing over signed α makes the sign inside ρ_{ij} immaterial.

Proof of Proposition E.1. Let $\hat{A} = \hat{\tilde{H}}_\Phi$ have full column rank and let $A^* = \tilde{H}_\Phi^*$ with $\|\hat{A} - A^*\|_2 \leq \delta_H$. The data obey $\tilde{Y} = A^* \mathbf{c} + \tilde{E} = \hat{A} \mathbf{c} + [\tilde{E} + (A^* - \hat{A}) \mathbf{c}]$. The least-squares estimate using $\hat{A}$ is $\hat{\mathbf{c}} = \hat{A}^\dagger \tilde{Y}$, and $\hat{A}^\dagger \hat{A} = I_J$, so

$$\hat{\mathbf{c}} - \mathbf{c} = \hat{A}^\dagger [\tilde{E} + (A^* - \hat{A}) \mathbf{c}].$$

Taking norms with $\|\hat{A}^\dagger\|_2 = 1/\sigma_J(\hat{A})$ and $\|(A^* - \hat{A}) \mathbf{c}\|_2 \leq \delta_H \|\mathbf{c}\|_2$ gives

$$\|\hat{\mathbf{c}} - \mathbf{c}\|_2 \leq \frac{\|\tilde{E}\|_2 + \delta_H \|\mathbf{c}\|_2}{\sigma_J(\hat{A})},$$

which is Equation 18.

Table 1: Identifiability diagnostics and reporting outputs produced with source activity estimates. The lower block lists the set- and graph-valued outputs that accompany the scalar diagnostics.

Symbol	Name	Interpretation	Action
$r_{\text{num}}(\tilde{H}_\Phi)$	numerical rank	exact separability	require rank J
$\sigma_J(\tilde{H}_\Phi)$	padded min. singular value	noise robustness	flag zero/small
$\kappa(\tilde{H}_\Phi)$	condition number	error amplification	∞ if deficient
r_{eff}	effective rank	noise-level resolution	reduce grouping
v_j	coefficient visibility	measurable effect	flag weak
a_j	background absorption	signal lost to Q	warn after projection
ρ_{ij}	pairwise coherence	pairwise ambiguity	merge coherent
d_{ij}^{ray}	ray distance	ray separation	N/A if weak
$\mathcal{W}$	weak set	undetectable coefficients	flag, no edges
$\mathcal{A}$	ambiguous-pair set	confusable pairs	source edges
$\mathcal{E}_{\text{src}}$	source ambiguity graph	source-level ambiguity	retain triggers
$\mathcal{G}$	connected report groups	identifiable resolution	report grouped
n/a	global unresolved warning	deficient, no edge	warn globally

B Identifiability Diagnostics: Definitions

This appendix gives the full definitions of the diagnostics summarized in Section 4 and Table 1.

For floating-point computation, the default numerical rank tolerance is $\tau_{\text{num}} = \max(N, J) \epsilon_{\text{mach}} \sigma_1$ with numerical rank $r_{\text{num}} = \#\{i : \sigma_i > \tau_{\text{num}}\}$, where ϵ_{mach} is the machine precision of the SVD dtype (about 2.2×10^{-16} in double precision). This τ_{num} is a purely numerical tolerance for detecting exact rank deficiency; it is distinct from the scientific, noise-dependent threshold τ_σ below. The primary identifiability score is the *padded minimum singular value* $\sigma_J(\tilde{H}_\Phi)$, where the singular spectrum is padded with zeros whenever $J > N$ or $\tilde{H}_\Phi$ is rank deficient. A rank-deficient response therefore scores zero and is never reported as stable. The condition number is $\kappa(\tilde{H}_\Phi) = \sigma_1/\sigma_J$ when $r_{\text{num}} = J$ and ∞ when $r_{\text{num}} < J$. The smallest numerically positive value may additionally be reported as $\sigma_{\text{min},+}$, but it is not the identifiability score and never turns a rank-deficient matrix into a stable one. For a separately chosen, noise-dependent threshold τ_σ , the effective rank is $r_{\text{eff}}(\tau_\sigma) = \#\{i : \sigma_i(\tilde{H}_\Phi) > \tau_\sigma\}$. Unlike τ_{num} , the threshold τ_σ is set from the observation-noise scale, not from machine precision. Given a per-row noise standard deviation σ_e (from sensor calibration or a bootstrap residual estimate), a singular direction is only resolvable when its singular value exceeds the level at which iid noise projects

onto it; a workable predeclared default is $\tau_\sigma = \sigma_e \sqrt{N}$. It is recorded independently of τ_{num} and is not tuned on source-recovery results. If $r_{\text{eff}}(\tau_\sigma) < J$, the sensor network does not support reliable estimation of all J source-basis coefficients at that noise level.

For coefficient-specific diagnostics, let $\tilde{\mathbf{h}}_j$, $j = (k, b)$, be a projected source-basis fingerprint. Its visibility and the weak set at threshold $\tau_v \geq 0$ are $v_j = \|\tilde{\mathbf{h}}_j\|_2$ and $\mathcal{W} = \{j : v_j \leq \tau_v\}$. A weak coefficient has little measurable effect after transport and background correction: v_j is the sensor-time magnitude produced by one unit of coefficient c_j , and τ_v marks fingerprints below the minimum detectable signal. A practical predeclared rule sets τ_v from the noise scale and a target signal-to-noise ratio, e.g. $\tau_v = \text{SNR}_{\text{min}} \sigma_e$, so a unit coefficient is flagged weak when even its full response is indistinguishable from noise. For eligible indices $i, j \notin \mathcal{W}$, pairwise ambiguity is measured by projected fingerprint coherence $\rho_{ij} = |\tilde{\mathbf{h}}_i^\top \tilde{\mathbf{h}}_j| / (\|\tilde{\mathbf{h}}_i\|_2 \|\tilde{\mathbf{h}}_j\|_2)$. Values near one indicate nearly proportional sensor-time signatures. The ambiguous-pair set is $\mathcal{A} = \{(i, j) : i, j \notin \mathcal{W}, \rho_{ij} > \tau_\rho\}$, $\tau_\rho \in [0, 1]$. The threshold τ_ρ is predeclared or calibrated on controlled experiments, not optimized after inspecting source recovery; a conservative default takes τ_ρ close to one so that only near-proportional fingerprints are merged. Coherence and ray distance are reported only for eligible (nonweak) pairs; entries involving a weak fingerprint are shown as “N/A” and excluded from extrema, while the weak-visibility flag is always reported.

To express the same scale-invariant geometry as a distance, we also report the ray distance for eligible, generally signed projected fingerprints, $d_{ij}^{\text{ray}} = \min_{\alpha \in \mathbb{R}} \|\tilde{\mathbf{h}}_i - \alpha \tilde{\mathbf{h}}_j\|_2 / \|\tilde{\mathbf{h}}_i\|_2$. For nonweak fingerprints,

$$d_{ij}^{\text{ray}} = \sqrt{1 - \rho_{ij}^2}, \quad (9)$$

which we derive in Appendix A; numerical implementations evaluate $\sqrt{\max(0, 1 - \rho_{ij}^2)}$. This is an unoriented linear-ray diagnostic, not a claim that projected fingerprints are non-negative. High coherence and small ray distance encode the same pairwise geometry, so ray distance is a diagnostic geometrically equivalent to coherence under this normalization, not an independent identifiability criterion. Like coherence, it is “N/A” for pairs involving $\mathcal{W}$ and excluded from distance extrema.

Background absorption. Background correction can improve robustness by removing broad temporal variation, but it can also remove source-driven signal. Let $P_Q = QQ^\dagger$. For coefficient fingerprint $j = (k, b)$, define the background absorption fraction $a_j = \|P_Q \mathbf{h}_j^{\text{lag}}\|_2 / \|\mathbf{h}_j^{\text{lag}}\|_2$. If $\|\mathbf{h}_j^{\text{lag}}\|_2$ is zero at numerical tolerance, a_j is undefined and reported as N/A; the coefficient remains explicitly weak. If $a_j \approx 1$, most of the raw source-basis fingerprint lies in the background space, and the projected visibility v_j will be small even if the raw fingerprint is large. This is why identifiability must be assessed after projection.

Table 2: Indices and objects in the temporal-basis apportionment model.

Symbol	Shape or range	Meaning
k	$1, \dots, K$	source-group index
b	$1, \dots, B$	temporal-basis index
$j = (k, b)$	$1, \dots, J$	source-major, basis-minor coefficient index
Φ	$T \times B$	nonnegative source-activity basis
C	$K \times B$	nonnegative source-basis coefficients
$\mathbf{c}$	$J = KB$	source-major vectorization of C
H_Φ^{lag}	$N \times J$	observed-row coefficient fingerprints
$\tilde{H}_\Phi$	$N \times J$	background-projected fingerprints
$\tilde{\mathbf{h}}_{kb}$	N	fingerprint for coefficient c_{kb}

C Temporal-Basis Notation and Diagnostic Conventions

The constant-activity model uses $B = 1$, $\phi_1(t) = 1$, and $c_{k1} = \theta_k$. In the general model, diagnostics are computed for the J coefficient fingerprints. Source-level activity trajectories are reconstructed as

$$\hat{\theta}_k(t) = \sum_{b=1}^B \hat{c}_{kb} \phi_b(t);$$

fingerprint columns are not summed to create an artificial source-level diagnostic vector.

For a source-level ambiguity summary, IASA reports the largest eligible coherence between coefficients belonging to different sources and retains the source-basis pair that attains it. Coherence and ray distance are undefined for pairs involving a coefficient with $v_j \leq \tau_v$; tables display these entries as N/A, exclude them from pairwise extrema, and report the weak coefficient separately.

D Transport Response Construction and Fit Details

This appendix gives the construction, background, and estimation details deferred from Section 3.

Wind imputer query and unit conversion. The observed station vectors and masks are supplied to a coordinate-query wind imputer: given sparse station observations $\{(\mathbf{z}_i, t, \mathbf{u}_{i,t}, m_{i,t})\}$, it predicts transport vectors at arbitrary spatial query locations and times. For the IASA response operator, we query the imputer on every response-grid cell $\mathbf{x}_g$ and hour t , producing

$$\hat{\mathbf{w}}_t(\mathbf{x}_g) = \hat{f}_\omega(\mathbf{x}_g, t; \{(\mathbf{z}_i, t', \mathbf{u}_{i,t'}, m_{i,t'})\}), \quad g = 1, \dots, n. \quad (10)$$

This yields a gridded wind field $\hat{\mathbf{W}} \in \mathbb{R}^{T \times n \times 2}$, rather than a station-only or city-averaged sequence. Real-data validation is performed by masking observed station vectors and measuring held-out station-time error; dense city-wide wind truth is not assumed. Meteorological speed must be converted to grid

displacement before puff advection. For response timestep Δt_s seconds and cell dimensions $\Delta x_m, \Delta y_m$, we use

$$v_x^{\text{grid}}(\mathbf{x}_g, t) = \frac{\hat{U}_x(\mathbf{x}_g, t) \Delta t_s}{\Delta x_m}, \quad v_y^{\text{grid}}(\mathbf{x}_g, t) = \frac{\hat{V}_y(\mathbf{x}_g, t) \Delta t_s}{\Delta y_m}. \quad (11)$$

Every response artifact records these scales, the coordinate convention, the wind units, and the exact converted sequence.

Wind-field ensembles. The imputed gridded wind field is an estimated transport input, not ground truth. To propagate wind-imputation uncertainty, we construct wind-field ensembles $\{\hat{\mathbf{w}}_t^{(r)}(\mathbf{x}_g)\}_{r=1}^R$ using held-out-calibrated prediction error, station bootstrap resampling, checkpoint ensembles, or perturbations matched to validation residuals. Each ensemble member is queried on the same response grid and converted to grid displacement before response construction. The resulting response matrices are tagged as transport ensembles and are kept separate from inventory-robustness scenarios.

Puff dynamics and dispersion. We instantiate the response with a Gaussian puff approximation. For a unit release from source cell i at location $\mathbf{r}_i$ and release time τ , the puff center is initialized as $\mathbf{z}_i(\tau) = \mathbf{r}_i$ and advected by the interpolated wind field, $d\mathbf{z}_i(a)/da = \hat{\mathbf{w}}_a(\mathbf{z}_i(a))$, $a \in [\tau, t]$. Source cells have integer centers and domain extents $[-0.5, N_x - 0.5] \times [-0.5, N_y - 0.5]$. With substep size δa , the trajectory is discretized as $\mathbf{z}_i^{r+1} = \mathbf{z}_i^r + \delta a \hat{\mathbf{w}}_{a_r}(\mathbf{z}_i^r)$. The final substep is shortened so that the trajectory lands exactly at the requested observation time. If the center exits Ω , the whole remaining puff is removed and never reflected, wrapped, clamped, or reinserted. Otherwise its contribution is spread with an anisotropic Gaussian kernel

$$K_\psi(\mathbf{x}; \mathbf{z}_i, \Sigma_i) = \frac{\exp\left(-\frac{1}{2}\|\mathbf{x} - \mathbf{z}_i\|_{\Sigma_i^{-1}}^2\right)}{2\pi|\Sigma_i|^{1/2}}, \quad (12)$$

where ψ denotes dispersion parameters. A simple covariance model is

$$\Sigma_i(t, \tau) = R_i(t, \tau) \begin{bmatrix} \sigma_{\parallel}^2(t - \tau) & 0 \\ 0 & \sigma_{\perp}^2(t - \tau) \end{bmatrix} R_i(t, \tau)^\top, \quad (13)$$

with R_i aligned to the mean wind direction along the trajectory. When the mean wind norm is below the numerical threshold, the downwind direction is undefined; the implementation uses the positive- x axis as a deterministic tie-breaker for the anisotropic covariance orientation, recorded in the response metadata and affecting only near-calm releases. The age in this covariance is lower-bounded by a positive minimum dispersion time, so same-time releases remain finite: effective age = $\max(t - \tau, t_{\min})$.

Sensor evaluation and loss bookkeeping. Response-matrix entries are computed by evaluating the kernel directly at sensor coordinates:

$$[OG_{t,\tau}^\partial(\hat{\mathbf{w}}_{\tau:t}; \psi)\mathbf{e}_i]_s = \chi_i(t, \tau) K_\psi(\mathbf{x}_s; \mathbf{z}_i(t), \Sigma_i(t, \tau)). \quad (14)$$

Here $\chi_i(t, \tau)$ is an open-boundary survival indicator for a release from source cell i at release time τ : it is one while the advected kernel center remains inside the modeled domain and zero after the release has exited. Boundary and truncation

losses are recorded only as diagnostic quantities and are never used to renormalize the sensor response. Boundary loss denotes the portion of the Gaussian kernel lying outside the modeled spatial domain, while truncation loss denotes the portion omitted by evaluating retained mass only over a finite kernel support; both are estimated separately by grid quadrature over the response grid. These kernel losses are distinct from whole-release exit loss, where the advected release center leaves the open domain and no later sensor contribution is assigned to that release. Because the response matrix evaluates concentration at sensor points rather than integrating over area-weighted grid cells, summing entries across sensors or repeated sensor-time observations does not estimate retained pollutant mass.

Baseline policy and differentiability. Response matrices are constructed for source-induced increments relative to a declared initial-condition baseline. The baseline policy is fixed before source fitting, recorded with the response artifact, and applied consistently to the observations and response rows. Application-specific choices, such as whether to subtract a zero-source transported initial field or to use a zero-initial-state response, are part of the experimental instantiation rather than the generic response definition. The open-boundary lagged plume response is implemented in PyTorch; tensor operations retained in the graph can be differentiated, but discrete exit events and reporting logic are not assigned gradients.

Background basis and implicit projection. The normal background basis may contain a global constant, centered and scaled elapsed-time polynomials, local-clock daily sine/cosine harmonics, reference-coded day effects, standardized regional sensor-coordinate trends, reference-coded sensor offsets, and a row-aligned user basis. It depends only on timestamps, day labels, sensor identity, and sensor coordinates, and its effective rank is capped at eight. Once the primary Q is fixed, its coefficients are estimated implicitly by projection, or explicitly in the joint fit; each alternative source-independent basis reports the resulting visibility, absorption, rank, conditioning, and ambiguity components. Letting $Q = U\Sigma V^\top$ be a thin singular value decomposition and retaining columns U_r above the recorded rank tolerance, projection is applied implicitly,

$$P_Q^\perp X = X - U_r(U_r^\top X), \quad (15)$$

without constructing an $N \times N$ matrix. The row metadata for Y , H_Φ^{lag} , and Q must be identical: each row must refer to the same timestamp, sensor identity, and sensor index in the same order before projection is applied.

Regularizers and joint fit. The regularizer in Equation 6 can be ridge stabilization $R(\mathbf{c}) = \|\mathbf{c}\|_2^2$ or a weak prior around inventory-derived activity levels $R(\mathbf{c}) = \|\mathbf{c} - \mathbf{c}_0\|_2^2$. Equivalently, one may fit background and source coefficients jointly,

$$(\hat{\mathbf{c}}, \hat{\gamma}) = \arg \min_{\mathbf{c} \geq 0, \gamma} \|Y - H_\Phi^{\text{lag}} \mathbf{c} - Q\gamma\|_2^2 + \lambda R(\mathbf{c}). \quad (16)$$

The projected formulation is used for identifiability because it exposes the part of each source–basis fingerprint that remains after background correction.

Constrained end-to-end refinement. IASA first constructs a fixed response matrix from the pretrained wind field and

default dispersion parameters, fits source activities, and computes identifiability diagnostics for that declared response. It then optionally performs a constrained local refinement of the wind field, dispersion parameters, source coefficients, and background coefficients, initialized at the fixed-response solution and allowed only to make small physically constrained corrections, so that improved sensor fit does not come from arbitrary changes to transport. Let ϕ denote the wind-imputer parameters and ψ the dispersion parameters of the puff response (such as $\sigma_\parallel$, $\sigma_\perp$, and the minimum dispersion age), with ϕ_0 and ψ_0 their pretrained or default values. We use R_{sm} for a predeclared wind-field smoothness penalty (temporal or spatial squared differences) and Ψ_{phys} for the physically admissible dispersion-parameter set. The refinement objective is

$$\begin{aligned} \mathcal{L}_{\text{refine}} = & \|Y - H_\Phi^{\text{lag}}(\phi, \psi) \mathbf{c} - Q\gamma\|_2^2 + \lambda_\theta R(\mathbf{c}) \\ & + \lambda_w \|\hat{\mathbf{w}}_{1:T}^\phi - \hat{\mathbf{w}}_{1:T}^{\phi_0}\|_2^2 + \lambda_\psi \|\psi - \psi_0\|_2^2 \\ & + \lambda_{\text{sm}} R_{\text{sm}}(\hat{\mathbf{w}}^\phi), \end{aligned} \quad (17)$$

subject to $\mathbf{c} \geq 0$, $\psi \in \Psi_{\text{phys}}$, and $\|\hat{\mathbf{w}}_{1:T}^\phi - \hat{\mathbf{w}}_{1:T}^{\phi_0}\|_\infty \leq \epsilon_w$. The refined response is accepted only as a constrained local correction (acceptance criteria in Appendix G) and is reported with its own response diagnostics. Without these restrictions, wind and plume parameters could fit sparse pollution sensors while weakening the physical and source-level interpretation of $\hat{\mathbf{c}}$ and the reconstructed activities.

E Response-Matrix Perturbations

This appendix expands the response-error discussion of Section 4. Write the projected response discrepancy as $\Delta_H = \Delta_{\text{tr}} + \Delta_{\text{inv}} + \Delta_{\text{int}}$, where Δ_{tr} is the transport term, Δ_{inv} the inventory term, and Δ_{int} their interaction. The decomposition is used to keep uncertainty honest: Δ_{tr} is propagated through transport ensembles and may become an uncertainty interval, whereas Δ_{inv} is reported only as named robustness scenarios and is never pooled into an interval, and Δ_{int} is recorded rather than pooled separately. Let $\hat{H}_\Phi$ be the true projected response and $\tilde{H}_\Phi$ the constructed one, with $\|\hat{H}_\Phi - \tilde{H}_\Phi^*\|_2 \leq \delta_H$. Then $\tilde{Y} = \hat{H}_\Phi \mathbf{c} + [\tilde{E} + (\tilde{H}_\Phi^* - \hat{H}_\Phi) \mathbf{c}]$, so response error acts like additional observation error.

Proposition E.1 (Perturbation bound). *Assume $\hat{H}_\Phi$ has full column rank and $\|\hat{H}_\Phi - \tilde{H}_\Phi^*\|_2 \leq \delta_H$. The unconstrained least-squares estimate using $\tilde{H}_\Phi$ satisfies*

$$\|\hat{\mathbf{c}} - \mathbf{c}\|_2 \leq \frac{\|\tilde{E}\|_2 + \delta_H \|\mathbf{c}\|_2}{\sigma_J(\hat{H}_\Phi)}. \quad (18)$$

A short proof, reusing the pseudoinverse bound of Proposition 4.3, appears in Appendix A. The same norm bound applies to all three discrepancy terms, but their interpretations are not interchangeable. Wind and dispersion ensembles quantify transport uncertainty and may support uncertainty intervals under the ensemble model. Alternative inventory maps, locations, labels, or scales are inventory robustness

Table 3: IASA thresholds, their meaning, and how they are set. None is optimized on source-recovery error.

Threshold	Meaning	How set
τ_{num}	numerical rank tol.	$\max(N, J)\epsilon_{\text{mach}}\sigma_1$
τ_σ	effective-rank floor	noise scale, e.g. $\sigma_e\sqrt{N}$
τ_v	visibility floor	$\text{SNR}_{\min}\sigma_e$
τ_ρ	ambiguity coherence	predeclared, near 1
τ_ρ^{ref}	refinement coherence	default $\tau_\rho^{\text{ref}} = \tau_\rho$
ϵ_w	wind-drift cap	physical wind tolerance

scenarios, not confidence intervals. They are reported separately, and source-specific attribution remains conditional on the supplied inventory. Either error type is most damaging when the projected response is already ill-conditioned.

F Per-Sensor Source Footprints

This appendix expands Section 4. For an observed row (s, t) , the background-corrected fit decomposes exactly as $\tilde{y}_{s,t} = \sum_{k,b} \tilde{H}_{\Phi,(s,t),(k,b)} \hat{c}_{kb}$, so the contribution of source k to sensor s over the record is $\hat{Y}_k^{(s)} = \sum_t \sum_b H_{\Phi,(s,t),(k,b)}^{\text{lag}} \hat{c}_{kb}$. The projected form isolates the identifiable part; the unprojected form gives the raw fitted sensor contribution before background removal. The spatial origin of a sensor’s signal is the pullback of its response row onto the source grid: for sensor s at time t , $F_{s,t}(i) = \sum_{\ell \in \mathcal{L}_t} [OG_{t,t-\ell}^{\partial}(\hat{\mathbf{w}}_{t-\ell:t})]_{s,i} \geq 0$ and $F_s(i) = \sum_t F_{s,t}(i)$, a nonnegative field over source cells i recording which upwind cells influence sensor s . Weighting by $\phi_b(t-\ell)$, the inventory map $\mathbf{s}_k$, and $\hat{c}_{kb}$ resolves the fitted contribution of each source group to each sensor by cell of origin. This is the puff response of Section 3 read backward from the sensor and requires no new transport operator. Per-sensor attribution uses only the rows of $\tilde{H}_\Phi$ for that sensor: for the row submatrix $\tilde{H}_\Phi^{(s)}$, $\tilde{H}_\Phi^{(s)\top} \tilde{H}_\Phi^{(s)} \preceq \tilde{H}_\Phi^\top \tilde{H}_\Phi$, so deleting rows cannot increase the smallest singular value or the rank and a single sensor is never more identifiable than the pooled network. Footprints and per-sensor contributions are therefore reported as spatial overlays of the global fit; per-sensor source *shares* are asserted only at the globally identifiable resolution, aggregating to a report group wherever the pooled diagnostics require one, which keeps the spatial view from implying separations the sensing system does not support.

G Uncertainty, Adequacy, and Reporting Details

This appendix expands the reporting components summarized in Section 4, and gives the IASA thresholds (Table 3).

Uncertainty reporting. Let $\tilde{r} = \tilde{Y} - \tilde{H}_\Phi \hat{\mathbf{c}}$ be the projected residual and $\hat{\sigma}^2 = \|\tilde{r}\|_2^2 / \max(N - r_{\text{eff}}(\tau_\sigma), 1)$ a residual variance estimate. For an unconstrained or active-set approximation, let $\mathcal{A}_c = \{j : \hat{c}_j > \tau_c\}$ and let $\tilde{H}_{\Phi,\mathcal{A}_c}$ contain the active columns. The covariance of the active estimates is approximated by

$$\text{Cov}(\hat{\mathbf{c}}_{\mathcal{A}_c}) = \hat{\sigma}^2 (\tilde{H}_{\Phi,\mathcal{A}_c}^\top \tilde{H}_{\Phi,\mathcal{A}_c} + \lambda I)^\dagger. \quad (19)$$

The use of $N - r_{\text{eff}}(\tau_\sigma)$ is a heuristic degrees-of-freedom count that does not exactly account for the background rank r or the NNLS active-set constraints. Equation 19 is the standard active-set / ridge least-squares covariance approximation (Seber and Lee 2003): its diagonal entries give per-coefficient uncertainty and its off-diagonal entries expose coupled (jointly uncertain) estimates. For each transport ensemble member from Section 3, IASA rebuilds $\tilde{H}_\Phi^{(b)}$, refits $\hat{\mathbf{c}}^{(b)}$, and reports empirical quantiles across members as transport uncertainty. Every ensemble is tagged by provenance: wind and dispersion members have `ensemble_kind=transport` and may form transport-uncertainty intervals, whereas alternative inventory maps, source locations, labels, or scales have `ensemble_kind=inventory` and produce scenario-wise robustness tables. The two kinds are never pooled into one interval, and active-set intervals describe the fitted observation model without converting inventory scenarios into probabilistic uncertainty.

Residual model-adequacy check. Let Z be an orthonormal basis for the background-orthogonal observed-row space, and define the nondegenerate residual coordinates and externally supplied noise covariance

$$\bar{r} = Z^\top (Y - H_\Phi^{\text{lag}} \hat{\mathbf{c}}), \quad \bar{\Sigma}_e = Z^\top \Sigma_e Z, \quad (20)$$

where Σ_e is an externally declared observation-noise covariance and Z removes the background directions so that $\bar{r}$ carries no degrees of freedom absorbed by Q . For independent homoscedastic noise $\Sigma_e = \sigma_e^2 I$, orthonormality of Z gives $\bar{\Sigma}_e = \sigma_e^2 I$ and $T_{\text{res}} = \|\bar{r}\|_2^2 / \sigma_e^2$. When $\bar{\Sigma}_e$ is calibrated independently of this fitted residual and is positive definite, the adequacy statistic is

$$T_{\text{res}} = \bar{r}^\top \bar{\Sigma}_e^{-1} \bar{r}. \quad (21)$$

Its null distribution is not a plain χ^2 , because $\hat{\mathbf{c}}$ is fitted from the same data; IASA therefore calibrates it by parametric bootstrap. It simulates $B_{\text{res}} = 1000$ synthetic observation sets from the fitted source/background model plus the declared noise model, refits the coefficients on each draw, and recomputes T_{res} to form the null distribution. At $\alpha = 0.05$ it flags inadequacy when the observed statistic exceeds the empirical $(1 - \alpha)$ -quantile; the Monte Carlo p -value uses the standard add-one correction. If no externally calibrated noise model is supplied, IASA reports raw and projected residual norms and sensor-wise, time-wise, and autocorrelation summaries with `calibration_status=uncalibrated` and emits no pass/fail adequacy claim. Rejection establishes a mismatch with the declared model but does not identify a missing source, and non-rejection does not establish inventory completeness: an omitted source whose sensor-time signature lies in $\text{span}[H_\Phi^{\text{lag}}, Q]$ can be absorbed by fitted terms and remain residual-invisible.

Wind-distribution diagnostics. Diagnostics for one realized wind window answer a retrospective question. Prospective network adequacy is assessed empirically over predeclared historical or simulated wind windows. Across those response matrices, IASA reports the 5th, 50th, and 95th percentiles of σ_J , the probability of full numerical and effective rank, the probability that each coefficient is weak, the probability

Table 4: New Delhi source groups, spatial proxies, and temporal-basis components.

Source group	Spatial proxy	Temporal-basis components
Brick kilns	regional kiln map	alternating 12-hour blocks
Industries	regional industry map	continuous + daytime (07–19)
Population density	gridded population	peaks 07:00, 13:00, 19:00
Traffic	road-network map	diurnal slots 00/06/12/18

that each source pair is ambiguous, and the frequency of each conservative report component. These are Monte Carlo or empirical distributional summaries under the sampled wind population, not a new identifiability theorem.

Acceptance criteria for refinement. If the optional constrained refinement of Appendix D is used, it is accepted only if it improves fit without degrading separability. Let $\tilde{H}_{\Phi,0}$ be the projected response before refinement and $\tilde{H}_{\Phi,\text{ref}}$ after refinement. We require $\sigma_J(\tilde{H}_{\Phi,\text{ref}}) \geq (1 - \eta_{\text{id}})\sigma_J(\tilde{H}_{\Phi,0})$ and $\max_{i \neq j, i,j \notin \mathcal{W}} \rho_{ij}^{\text{ref}} \leq \tau_{\rho}^{\text{ref}}$. These checks prevent end-to-end tuning from improving sensor fit by making source-basis fingerprints less distinguishable. If either response has fewer than J numerically nonzero singular values, its σ_J is taken to be zero (Section 4), so refinement cannot be accepted by exploiting a rank-deficient response.

Reporting protocol. The source ambiguity graph, its deterministic connected components, and the transitive over-merging of an A – B – C chain are defined in Section 4. Operationally, those components are the conservative recommended report groups, each edge retains its triggering coefficient pair, and rank deficiency without an eligible pair yields a global unresolved warning rather than an invented edge. For each reported source group G_a , IASA reports the activity trajectory $\hat{\theta}_{G_a,1:T}$, its interval $\text{CI}(\hat{\theta}_{G_a})$, member visibilities $\{v_{kb} : k \in G_a\}$, and the maximum eligible cross-group coherence $\max_{k \in G_a, k' \notin G_a, (k,b), (k',b') \notin \mathcal{W}} \rho_{(k,b),(k',b')}$. A source contribution is marked reliable only when its constituent reported coefficients are above the visibility threshold, its uncertainty interval is sufficiently narrow, and its eligible coherence with other reported groups is below the ambiguity threshold. Reports retain the basis names of every weak coefficient and the coefficient pair attaining each cross-source coherence maximum; coherence and ray distance involving weak fingerprints are displayed as “N/A” and excluded from extrema, while the weak flags remain visible. Otherwise IASA labels the source as weakly visible, ambiguous with another source group, or merged into a coarser identifiable group.

H New Delhi Platform Details

This appendix expands the platform construction of Section 5: the source inventories (Table 4), the kriged initial-condition baseline, and the auxiliary simulator.

Observations, sensors, and inventory preparation. The

hourly government records (CPCB 2023, 2025) (columns `monitor_id`, `timestamp_round`, `AT`, `RH`, `WD`, `WS`, `pm10`, `pm25`) span 21,960 hourly timestamps from 1 May 2018 through 31 October 2020 in Indian Standard Time; approximately 74.9% of wind-direction, 72.8% of wind-speed, and 90.5% of $\text{PM}_{2.5}$ entries are observed, with a valid-value mask retained alongside station coordinates and timestamps. The two Pusa monitors (`Pusa_IMD` and `Pusa_DPCC`) are averaged by timestamp into `Pusa_averaged` (mean coordinates, per-variable hourly averaging), giving the final 32-sensor layout; $\text{PM}_{2.5}$ is never imputed and its flattened valid-value mask defines $M_{\mathcal{O}}$, with the same ordered sensor-time rows retained in Y , H_{Φ}^{lag} , Q , and metadata. Brick-kiln and industry proxies follow the regional emissions inventory (Guttikunda and Calori 2013), population density the gridded population product (Center for International Earth Science Information Network (CIESIN), Columbia University 2018), and the traffic map is derived from road map data (Google 2024). Every 80×80 map is cropped to the New Delhi study window (rows 21:61, columns 16:56) covering the regulatory network, giving a 40×40 map, then divided by its own cropped 99th percentile. Traffic is modeled as a single road-network source whose nearest-slot maps at 00:00, 06:00, 12:00, and 18:00 define traffic-specific temporal-basis components under the shared dictionary Φ and the fixed-zero mask $\mathcal{F}_0$ (assuming a time-invariant congestion spatial pattern whose amplitude varies); brick-kiln activity uses deterministic alternating 12-hour blocks, industry a continuous operating fraction with higher daytime (07:00–19:00) activity, and population a baseline with smooth cooking peaks. These are deterministic study assumptions, not observed emissions truth, and each map’s independent proxy normalization is why fitted coefficients and contributions are in normalized proxy units rather than physical emission totals or shares.

Forward-model baselines and study modes. Both forward paths use the New Delhi-derived 40×40 domain and regulatory sensor coordinates, with the open-boundary Gaussian puff operator evaluated directly at sensors. The same operator is the matched generator for controlled recovery experiments and the inference operator for observed $\text{PM}_{2.5}$; only the initial-condition baseline differs by mode: a zero-source, zero-initial-state baseline for controlled operator-matched runs, and a subtracted zero-source transported kriged-initial-condition baseline for observed runs. Normal backgrounds follow Section 3, with the base controlled background the rank-four member. Where a study is explicitly about inventory geometry (the coherence and inventory-robustness experiments) the real New Delhi inventory maps are used; the remaining controlled experiments use compact synthetic source cells for tractability, since their conclusions concern conditioning, coherence, and recovery rather than specific inventory footprints.

Kriged initial-condition baseline. For observed New Delhi runs, the initial pollution field is estimated from the first available regulatory observations (the two Pusa monitors averaged) by a normalized Gaussian-kernel spatial interpolation on the response grid, a kriging surrogate that avoids an additional heavyweight geostatistics dependency. Its zero-source transported trajectory, obtained by propagating this initial

Table 5: Experiment 1: coefficient recovery versus conditioning and noise. The better-conditioned geometry recovers more accurately at matched noise.

geom.	noise	σ_J	κ	coef. err.	resid.
separated	0.00	3.27	8.89	0.000	0.00
separated	0.05	3.27	8.89	0.066	6.37
separated	0.10	3.27	8.89	0.132	12.73
separated	0.20	3.27	8.89	0.264	25.46
close	0.00	9.60	2.26	0.000	0.00
close	0.05	9.60	2.26	0.028	8.42
close	0.10	9.60	2.26	0.057	16.84
close	0.20	9.60	2.26	0.113	33.68

field as a single release through the same open-boundary puff operator (its Gaussian kernel is the advection–diffusion Green’s function, so units are preserved), is subtracted from the sensor observations before fitting source activities. Source coefficients are therefore fitted to the residual sensor-time signal after removing the contribution of the initialized background field and the low-rank temporal background basis. Kernel interpolation from sparse first-hour observations is itself an ill-posed spatial estimate, so this baseline injects some uncertainty; the effect is limited by declaring and subtracting the baseline before fitting and by the low-rank background basis Q , which absorbs residual broad structure. Controlled operator-matched experiments may instead use a zero initial field when the synthetic data are generated from the same zero-initial-state response.

Auxiliary advection–diffusion simulator. The auxiliary advection–diffusion simulator uses first-order upwind advection, a five-point Laplacian, a two-stage Heun update, diffusivity 3×10^{-4} , and edge-hold boundaries. It provides a structural forward-model mismatch: data generated by this distinct transport operator are fit with the open-boundary puff response, testing the puff approximation against a different operator family. This mismatch also drives the residual-adequacy test of Section 4, letting us check its power against structural transport misspecification and not only against missing sources.

I Expanded Results

Full numeric tables backing the Section 5 figures: the seven controlled diagnostics and the baseline comparison of Figure 1, and the observed New Delhi weeks of Figure 2.

I.1 Conditioning (Exp 1)

Coefficient recovery versus conditioning and noise; the better-conditioned geometry recovers more accurately at matched noise.

I.2 Wind Diversity and Sensor Geometry (Exp 4)

σ_J , maximum eligible coherence, and coefficient error across wind provider $\times$ sensor layout, and the wind-window ensemble spread.

Table 6: Experiment 4: σ_J , maximum eligible coherence, and coefficient error across wind provider $\times$ sensor layout.

wind	layout	σ_J	max. coh.	coef. err.
constant	regulatory	1.53	0.975	0.000
constant	random	0.00	0.028	0.000
constant	downwind	0.13	0.990	0.000
single	regulatory	6.44	0.741	0.000
single	random	0.00	1.000	0.573
single	downwind	1.36	0.990	0.000
diurnal	regulatory	6.11	0.722	0.000
diurnal	random	0.00	0.880	0.001
ar1	regulatory	7.35	0.641	0.000
ar1	random	0.00	0.774	0.115
multi	regulatory	5.70	0.129	0.000
multi	random	2.79	0.322	0.000
real	regulatory	2.85	0.156	0.000
real	random	2.10	0.287	0.000
real	downwind	0.38	0.221	0.000

Table 7: Experiment 4: wind-window ensemble distribution (σ_J percentiles; both ensembles reach full rank with probability 1 and show no ambiguity).

ensemble	σ_J^5	σ_J^{50}	σ_J^{95}
historical (real slices)	1.45	4.51	9.96
simulated (AR(1))	3.50	5.58	6.80

I.3 Background Stress (Exp 3)

Redundant $\equiv$ primary (same span); the source-like stress basis collapses identifiability ($\sigma_J \rightarrow 0$, absorption $\rightarrow 1$).

I.4 Transport Error (Exp 5)

Parametric perturbations (wind speed, direction, dispersion); a structural PDE mismatch generated with the auxiliary edge-hold advection–diffusion solver yields adequacy rejection rate 1.0.

I.5 Inventory Robustness (Exp 6)

Location, scale, category, and map-version perturbations, each a separate scenario never pooled with transport error; recovery exact, σ_J tracks the version.

I.6 Lag-Window Selection (Exp 7)

Conditioning improves with lag; the selected $L = 16$ is chosen from Equation 3 alone, not fitted coefficients.

I.7 Temporal-Basis Recovery (Exp 9)

Activity-trajectory error stays below coefficient error at every noise level (fixed σ_J), a fixed geometric property of the temporal-basis map.

I.8 Baseline Comparison (Exp 11)

IASA against plain NNLS (identifiability layer ablated), chemical mass balance (unprojected), and a PMF/NMF receptor factorization, on the wind/geometry collapse and source-like background stress. IASA flags both and, under stress, recovers

Table 8: Experiment 3: background stress. Redundant $\equiv$ primary (same span); the source-like stress basis collapses identifiability ($\sigma_J \rightarrow 0$, absorption $\rightarrow 1$).

background	min. vis.	absorp.	σ_J	coef. err.
none	1.000	0.000	0.080	2.36
primary	0.972	0.235	0.080	2.61
redundant	0.972	0.235	0.080	2.61
stress	0.000	1.000	0.000	0.78

Table 9: Experiment 5: parametric transport error. Coefficient error rises with the operator error norm on the direction and dispersion axes (the speed axis is non-monotonic). Structural PDE mismatch: adequacy rejection rate 1.0.

axis	value	op. err.	σ_J	coef. err.
direction ($^\circ$)	0	0.00	0.87	0.00
direction ($^\circ$)	5	0.29	3.37	0.62
direction ($^\circ$)	10	0.60	7.07	0.80
direction ($^\circ$)	20	0.87	7.98	0.86
speed ($\times$)	1.25	0.29	2.09	0.45
speed ($\times$)	1.50	0.53	4.57	0.27
dispersion ($\times$)	1.50	0.19	1.42	0.50
dispersion ($\times$)	2.00	0.35	1.76	0.89

apportionment to < 0.01 share error; no baseline flags either case.

I.9 Observed New Delhi (weeks 1–4)

Per-week identifiability geometry, proxy apportionment (fraction of fitted sensor signal, not physical emissions), and sensor-fit/residual diagnostics (PM_{2.5} never imputed; adequacy uncalibrated).

J Additional Experiments

Three further controlled experiments not shown in the main-text figures.

J.1 Coherence and Grouped Reporting (Exp 2)

We form increasingly coherent source pairs by shifting a copy of one inventory map by a controlled spatial offset, so decreasing the shift drives coherence toward one. Shrinking the offset drives maximum eligible coherence from 0.05 to 0.96 while σ_J falls from 21.8 to 4.6, yet at every offset the conservative rule kept the two sources in separate singleton components with exact recovery (individual and grouped relative error 0; Table 17): high coherence alone did not force a merge here. The reporting layer retains *every* trigger edge without deduplication when a merge does fire (an A–B–C chain keeps both edges).

J.2 Missing-Source Adequacy (Exp 8)

On the real New Delhi platform with a non-empty rank-4 primary background, the refitted parametric bootstrap ($\alpha = 0.05$, 1000 replicates) is well-calibrated (null rejection rate 0.05) and has full power against a residual-visible omitted source

Table 10: Experiment 6: inventory robustness on the real maps. Recovery is exact in every scenario; σ_J tracks the inventory version.

scenario	σ_J	coef. err.
baseline	21.56	0.0
location shift	28.85	0.0
spatial rescale	47.00	0.0
alt. map version	46.01	0.0
category swap	21.56	0.0

Table 11: Experiment 7: lag-window selection. Conditioning improves with lag; the selected lag ($L = 16$) is chosen from geometry alone, not fitted coefficients.

lag L	σ_J	κ	#comp.	coef. err.
4	0.08	569.2	2	0.0
6	0.71	63.8	2	0.0
8	2.46	18.6	2	0.0
10	4.57	10.1	2	0.0
12	6.16	7.5	2	0.0
16	9.07	5.1	2	0.0

whose signature lies largely outside $\text{span}[H_\Phi^{\text{lag}}, Q]$ (rejection 1.0); an *aligned* omission, a nonnegative combination of the fitted columns and background, is absorbed and not detected (rejection 0.05; Table 18). The test is therefore one-sided: a rejection is evidence of model inadequacy, whereas a non-rejection cannot certify inventory completeness.

J.3 Per-Sensor Footprints (Exp 10)

On controlled trials the per-sensor footprints of Section 4 localize the responsible upwind cells (mass centroid 0.74 cells from the true origin, 100% of mass within the 4-cell radius, nonnegative) and per-sensor contributions sum to the fitted sensor signal to machine precision; because a non-singleton component is present, footprints are reported at group granularity. On observed New Delhi data footprints are reported per monitor at the identifiable report groups, accompanying the apportionment of Section 5.2 (Appendix F).

K Extended Discussion on Related Work

This appendix expands the four strands summarized in Section 2.

Air-pollution source apportionment. Source apportionment has a long history in air-quality science, with receptor-oriented methods estimating source contributions from ambient chemical or physical measurements and source-oriented methods propagating emissions inventories through atmospheric transport models. Classical receptor frameworks include chemical mass balance and factor-analytic approaches such as positive matrix factorization (Watson et al. 2002; Paatero and Tapper 1994; Hopke 2016). Reviews and harmonization efforts have clarified best practices for receptor models, source-oriented models, and their combined use, while also documenting practical sensitivity to source profiles, factor interpretation, chemical resolution, and model comparability (Viana et al.

Table 12: Experiment 9: temporal-basis recovery (bases: diurnal, block, day/night). Activity-trajectory error stays below coefficient error at every noise level.

noise	coef. err.	activity err.
0.00	0.000	0.000
0.02	0.132	0.075
0.05	0.329	0.188
0.10	0.659	0.375
0.20	0.900	0.582

Table 13: Experiment 11: IASA vs. identifiability-blind base-lines on two non-identifiable scenarios (seed 0). Apportionment share error and whether an identifiability flag was raised. IASA flags in both scenarios; no baseline does.

scenario	method	flag	share err.
collapse	IASA	yes	0.000
collapse	plain NNLS	no	0.000
collapse	CMB	no	1.414
collapse	PMF	no	0.085
bg. stress	IASA	yes	0.006
bg. stress	plain NNLS	no	1.231
bg. stress	CMB	no	0.189
bg. stress	PMF	no	1.152

2008; Belis et al. 2013, 2019; Mircea et al. 2020; Hopke et al. 2020). These sensitivities arise because ambient measurements provide only indirect mixtures of source contributions: source profiles may be incomplete or locally unavailable, tracers and proxy inventories can be collinear, factor decompositions may admit multiple plausible interpretations, and transport or inventory assumptions can map distinct source configurations to similar observations. This ambiguity is especially visible in settings such as India, where studies can produce divergent source-contribution estimates for the same city because of limited local profiles, collinear tracers, and methodological differences (Pant and Harrison 2012); global reviews further show that policy-relevant source categories such as traffic, industry, dust, and domestic combustion are routinely reported from heterogeneous apportionment studies (Karagulian et al. 2015). Our work is complementary to this literature: rather than proposing another receptor model or emissions inventory, we analyze when known or proxy source maps are distinguishable from sparse concentration sensors after wind-conditioned transport, finite lags, background correction, and noise.

Inverse modeling and physics-constrained learning. Estimating latent states or parameters from partial observations is central to inverse problems and data assimilation (Tarantola 2005; Engl, Hanke, and Neubauer 1996). Classical approaches such as Kalman filtering, ensemble Kalman filters, and variational data assimilation use dynamical models together with priors, covariance assumptions, and observation models to infer hidden quantities from measurements (Wang et al. 2023; Carrassi et al. 2018). More recent physics-informed and operator-learning methods, including PINNs

Table 14: Observed New Delhi identifiability, weeks 1–4. All weeks are full rank, finite-conditioned, with singleton report groups.

week	σ_1	σ_J	num. rank	max. coh.	groups
1	50.1	3.71	7	0.316	4 singletons
2	60.3	5.67	7	0.355	4 singletons
3	77.0	10.15	7	0.496	4 singletons
4	81.2	7.65	7	0.475	4 singletons

Table 15: Observed New Delhi proxy apportionment, weeks 1–4 (fraction of fitted sensor signal, not physical emissions). Note the large week-to-week swing.

week	brick kilns	industries	population	traffic
1	0.00	0.00	1.00	0.00
2	0.00	0.16	0.81	0.04
3	0.06	0.00	0.94	0.00
4	0.93	0.07	0.00	0.00

and neural operators, incorporate governing equations or physics-motivated structure into flexible learning systems for parameter estimation and system identification (Raissi, Perdikaris, and Karniadakis 2019; Li et al. 2024, 2020; Bhardwaj, Balashankar, and Subramanian 2025). These methods motivate our use of a transport-structured, PyTorch-based constrained fitting backend, but they do not by themselves resolve the identifiability question. In sparse sensing regimes, the map from latent source activities, transport parameters, and background components to sensor observations can be many-to-one: different parameterizations may fit the same measurements after optimization. Our focus is therefore complementary to calibration and physics-constrained learning methods. Rather than asking only how to fit the projected sensor data, we characterize when the fitted source-basis coefficients are identifiable from the projected lagged response matrix and when the result should be weakened, qualified, or grouped.

Sensor-space spatio-temporal learning. Graph neural networks, transformers, and sequence models have been widely used to model spatio-temporal processes from sensor data, including traffic, environmental, and air-quality time series (Li et al. 2017; Yu, Yin, and Zhu 2017; Wu et al. 2019; Chen et al. 2023; Iyer et al. 2022; Bhardwaj et al. 2025; Bhardwaj, Balashankar, and Subramanian 2025). These methods are relevant to our setting because they operate in the same sparse-sensor observational regime. However, their usual objective is forecasting, interpolation, or representation learning over sensor measurements rather than attribution to physically meaningful source inventories. A sensor-space model can predict observed concentrations well while leaving unresolved whether distinct source groups would produce distinguishable wind-conditioned fingerprints at the deployed sensors. Our work is therefore complementary: instead of learning an unconstrained predictive representation of the sensor field, we analyze the identifiability of declared source-basis contributions under a specified transport, background, lag, and

Table 16: Observed New Delhi sensor fit and residual diagnostics, weeks 1–4. $\text{PM}_{2.5}$ never imputed; adequacy uncalibrated (no verdict).

week	obs. rows	mask frac.	resid. norm	adequacy
1	4561	0.848	3768	uncalibrated
2	4278	0.796	5955	uncalibrated
3	4315	0.803	3050	uncalibrated
4	4236	0.788	3760	uncalibrated

Table 17: Experiment 2: coherence rises as the source offset shrinks, but the conservative rule merged no sources and recovery stayed exact.

offset	max. coh.	σ_J	merged	indiv./grp. err.
8.0	0.048	21.77	no	0.0 / 0.0
6.0	0.862	10.48	no	0.0 / 0.0
4.0	0.905	8.40	no	0.0 / 0.0
2.0	0.945	5.43	no	0.0 / 0.0
1.0	0.958	4.64	no	0.0 / 0.0

noise model.

Observability and system identification. The question of whether latent quantities can be recovered from observations is closely related to classical notions of observability in control theory (Chen 1984). For linear dynamical systems, observability characterizes whether the internal state can be uniquely determined from output trajectories. Extensions to nonlinear systems and PDEs have been studied in system identification and inverse-problem literature (Ljung 1987; Banks and Kunisch 2012; Colton et al. 1990), often under assumptions of sufficient measurement richness. We bring this perspective to inventory-based air-pollution source apportionment under sparse spatial sensing. Rather than asking whether a full latent state or physical model is observable, we ask whether declared source–basis contributions have distinct projected sensor-time fingerprints after wind-conditioned transport, finite lags, background correction, and noise. This leads to diagnostics tailored to the apportionment task, including rank and singular-value stability, source visibility, pairwise coherence, background absorption, and conservative grouping, together with a real-world instantiation that reports the attribution resolution supported by the available sensors, inventories, and transport assumptions.

L Implementation Invariants

The response is constructed first on the complete time-major, sensor-minor row index. The $\text{PM}_{2.5}$ observation mask then selects the same ordered rows from Y , H_Φ^{lag} , Q , and metadata. Source–basis columns are always source-major and basis-minor. Any row or column metadata mismatch is a hard error.

The response construction uses integer source-cell centers, half-cell domain extents, fractional final advection substeps, a positive minimum dispersion age, direct sensor evaluation, finite lag, and complete removal after center exit. Kernel boundary/truncation loss and whole-release exit loss are

Table 18: Experiment 8: missing-source adequacy on the platform (non-empty Q). Calibrated null, full power on a residual-visible omission, and no detection of an absorbed in-span omission.

omission case	rejection rate
none (null, calibration)	0.05
residual-visible (out of span, frac. 0.91)	1.00
aligned (in span, absorbed)	0.05

recorded separately. Retained and dropped kernel mass sum to emitted kernel mass within the recorded quadrature tolerance; sensor observations are never summed as a mass proxy. The default puff baseline has zero source and zero initial state.

Normal backgrounds have effective rank at most eight and do not use inventory or response columns. Thin-SVD projection stores the numerical tolerance, singular values, retained U_r , dependent input columns, removed components, and orthogonality/idempotence residuals. A source-like column is permitted only when the run is labeled as a stress test.

The primary Q , lag-candidate grid, and fixed-zero coefficient set are declared before fitting. A fixed-zero mask removes columns before diagnostics and fitting and stores both reduced-to-original and original-to-reduced index maps. Fitted near-zero coefficients never modify the mask. Adjacent lag responses use identical rows and columns; the default convergence threshold is $\tau_L = 10^{-3}$.

M Sanity Gates

Before paper experiments, the implementation must pass six small deterministic gates built from public APIs:

1. **S1: response.** On a 16×16 grid with four compact sources, four sensors, impulse/constant bases, and seed 123, verify response shape, downwind arrival, open-boundary exit, no wraparound, dispersion anisotropy, mass accounting, metadata, and exact repeatability.
2. **S2: projection.** Use the S1 response with a rank-four normal background to verify exact ordering, background removal, orthogonality, idempotence, redundant-column invariance, and the labeled source-like stress behavior.
3. **S3: diagnostics.** Verify orthogonal, duplicate, and weak matrices, including padded zeros, σ_J , infinite deficient condition status, weak-pair N/A values, and a matched multi-source eastward versus two-direction wind comparison.
4. **S4: fit.** Recover known nonnegative coefficients in noiseless/noisy well-conditioned cases, verify projected-FISTA KKT convergence, preserve exact nonnegativity, and recover the duplicate-pair sum when the individual split is nonunique.
5. **S5: report groups.** Verify the duplicate-source connected component, an A – B – C chain with both trigger edges retained, separated-source non-edge, independent weak flag, conservative-group flag, and summed group contribution.
6. **S6: end to end.** Build sources, wind, response, background, masked observations, diagnostics, fit, uncertainty, and

report groups with one command and write a compact provenance-complete summary.

Task 10 experiments begin only after S1–S6 and the PyTorch CPU/CUDA parity checks pass. Two further gates support the results. **S7 (adequacy calibration)** runs a statistically calibrated study of the residual-adequacy test on a controlled orthonormal design, checking that the null rejection rate matches α , that the null p -values are uniform, and that power increases monotonically with omission strength. A final **reporting gate** exercises the evaluation layer that emits the result tables in Section 5, verifying our reporting rules: undefined weak-pair metrics stay null, grouped metrics accompany any non-singleton report component, an A – B – C merge chain retains both trigger edges, percentages are labeled fractions of fitted sensor signal, and an uncalibrated run never reports an adequacy pass.

N PyTorch Solver and Uncertainty Protocol

After CSV/NPZ ingestion, numerical computation uses PyTorch tensors on an explicit device and dtype. The nonnegative coefficient objective is

$$f(\mathbf{c}) = \|\tilde{\mathbf{Y}} - \tilde{\mathbf{H}}_{\Phi}\mathbf{c}\|_2^2 + \lambda\|\mathbf{c} - \mathbf{c}_0\|_2^2, \quad \mathbf{c} \geq 0. \quad (22)$$

Projected FISTA uses a Lipschitz step derived from σ_1 or a recorded power-iteration estimate, clamps every iterate to the nonnegative orthant, and restarts acceleration whenever the objective increases. Termination requires the configured projected-gradient/KKT residual and relative objective-change tolerances, or reports that the iteration cap was reached. The fit artifact stores convergence status, iteration count, final KKT residual, objective summary, device, and dtype.

Active-set covariance uses PyTorch linear algebra on $\tilde{\mathbf{H}}_{\Phi, \mathcal{A}}^\top \tilde{\mathbf{H}}_{\Phi, \mathcal{A}} + \lambda \mathbf{I}$. Wind-ensemble and bootstrap refits are batched on the same device when shapes permit. Undefined weak-pair metrics are serialized as JSON `null`, never NaN or a false numerical distance.

Uncertainty artifacts distinguish ensemble_kind=transport from ensemble_kind=inventory. Only the former may be summarized as an ensemble uncertainty interval; inventory members remain named robustness scenarios, and an aggregation request that mixes the two kinds is an error.

O Adequacy and Ensemble Result Contracts

A residual-adequacy result stores `calibration_status`, T_{res} , bootstrap quantile, Monte Carlo p -value, α , replicate count, and noise-model provenance. Paper runs use $\alpha = 0.05$ and 1000 refitted replicates. If the externally calibrated noise covariance is absent or invalid, numerical test fields are null and the status is `uncalibrated`; no boolean pass is emitted. Sensor, time, and autocorrelation residual summaries are retained in either case.

A lag-selection result stores every candidate L , adjacent convergence ratio, selected L , τ_L , physical grid rationale, row count, and diagnostic/group summaries. A wind-distribution result stores the sampled-window provenance; 5th, 50th, and 95th percentiles of σ_J ; full numerical and effective-rank

probabilities; coefficient weakness probabilities; pairwise ambiguity probabilities; and conservative-component frequencies.

P Reproducibility and Artifact Provenance

Every paper run records the container image, repository revision, PyTorch and CUDA versions, device model, dtype, deterministic-algorithm settings, random seeds, source files and crop, per-source scales, wind-imputer architecture, held-out wind split, query grid, gridded wind product, and ensemble provenance, inventory identifier/hash/version, observation mask, response and transport configuration, wind realization, dispersion settings, background columns and rank tolerance, lag protocol, fixed-zero mask and index mapping, diagnostic thresholds, ensemble kind, solver settings, noise-model provenance, and parent artifact hashes. Generated wind products, response matrices, checkpoints, and experiment outputs are written to generated-data or log paths and are not treated as immutable source files.

Controlled-run artifacts contain true and estimated coefficients, reconstructed activities, response/projection results, diagnostics, uncertainty, report groups, and recovery metrics. Observed New Delhi artifacts omit unavailable ground truth and instead contain observed-row counts, fitted trajectories, raw and projected residuals, normalized proxy activities/contributions, uncertainty, weak/ambiguous flags, trigger pairs, and group summaries. A percentage is emitted only with an explicit denominator identifying it as a fraction of fitted inventory-attributed sensor signal.